

Backbone Determination in a Wireless Sensor Network

Jake Carlson

February 18, 2018

Abstract

A report on implementing algorithms to partition a random geometric graph into bipartite subgraphs. Three different graph geometries are explored: unit square, unit disk, and unit sphere. Nodes are uniformly distributed in the geometry. Then the edges are determined and the vertices are colored using smallest-last vertex ordering.

Contents

1	Executive Summary	3
1.1	Introduction	3
1.2	Environment Description	3
2	Reduction to Practice	3
2.1	Data Structure Design	3
2.2	Algorithm Descriptions	4
2.2.1	Node Placement	4
2.2.2	Edge Determination	4
2.2.3	Graph Coloring	4
2.3	Algorithm Engineering	5
2.3.1	Node Placement	5
2.3.2	Edge Determination	5
2.3.3	Graph Coloring	6
2.4	Verification	7
2.4.1	Node Placement	7
2.4.2	Edge Determination	7
3	Appendix A - Figures	9
4	Appendix B - Code Listings	25

Listings

1	Processing driver	25
2	Topology class and subclasses	25

Benchmark	Order	Avg Deg	Topology	r	Size	Realized Avg Deg	Max Deg	Min Deg	Max Deg Del
1	1000	32	Square	0.101	14397	28	45	8	19
2	8000	64	Square	0.050	245884	61	95	17	39
3	16000	32	Square	0.025	250059	31	51	8	25
4	64000	64	Square	0.018	2016843	63	97	10	40
5	64000	128	Square	0.025	4005101	125	181	30	75
6	128000	64	Square	0.013	4052365	63	97	12	42
7	128000	128	Square	0.018	8070473	126	179	38	73
8	8000	64	Disk	0.045	245420	61	89	20	38
9	64000	64	Disk	0.016	2021818	63	101	18	42
10	64000	128	Disk	0.022	4018364	125	180	48	74
11	16000	64	Sphere	0.126	513208	64	94	38	41
12	32000	128	Sphere	0.126	2048539	128	170	83	89
13	64000	128	Sphere	0.089	4095131	127	188	88	88

Table 1: Benchmarks for Generating and Coloring RGGs

1 Executive Summary

1.1 Introduction

Random geometric graphs (RGGs) are useful for simulating wireless sensor networks placed in different topologies. This project examines three different geometries: Square, Disk, and Sphere. The user supplies parameters for how many nodes they want in the network and how many connections they want for each node. Then, the simulation finds the average radius needed for that number of connections, finds multiple backbones for the RGG, and displays the results graphically.

1.2 Environment Description

The data structures and topologies for this simulation are implemented in Python2.7. The graphics are done using Processing.py. All development and benchmarking has been done on a 2014 MacBook Pro with a 3 GHz Intel Core i7 processor and 16 GB of DDR3 RAM running macOS High Sierra 10.13.3.

A separate data generation script was used to generate the graphs using matplotlib. This library, and a variety of others, could not be imported into Processing.py because the jython interpreter used by Processing only accepts libraries written in raw Python.

The different geometries were implemented in a stand alone Python file and imported into the Processing.py script or the data generation script depending on what was being run. Because there is no intermediary file to hold the generated nodes and edges, there is no additional disk space needed to run the simulation. Everything can be done in system memory managed by Processing.

2 Reduction to Practice

2.1 Data Structure Design

The primary data structure used for this project is an adjacency list. However, to allow for constant time lookup of edges of a node, I used a Python dictionary where the keys are nodes and the values are a list of adjacent nodes. The space needed by the adjacency list is $\Theta(2n)$ where $n = |E|$. Two entries are used for each edge because they are undirected. This is superior to the adjacency matrix data structure which would require $\Theta(n^2)$ space where $n = |E|$.

In order to make this project maintainable as it is developed along the semester, I used the object-oriented capabilities Python has to offer to design the different geometries. I start with a base Topology class that creates the interface Processing uses to draw the graphs. This base class implements all of the methods needed for node placement and edge detection in 2D graphs. Then, I create three subclasses: Square, Disk, and Sphere.

The Square and Disk topologies simply need to override the methods for generating nodes and calculating the node radius needed for the desired average degree. The Sphere subclass needs to override a few additional functions because it exists in a 3D space. Other than the methods for generating nodes and calculating the node radius, it also needs to override the function used to draw the graph so that Processing will render the graph properly in 3D.

2.2 Algorithm Descriptions

2.2.1 Node Placement

A different node placement algorithm is required for each of the geometries. For the Square, the coordinates for each node are generated as two random numbers taken from a uniform distribution on the range $[0, 1]$. All of these points are guaranteed to be in the unit square.

For the Disk, a similar method is used. The coordinates for nodes are randomly sampled from a uniform distribution; however, if a node has a distance from the center of the Disk greater than the radius of 1, the coordinates for that node are resampled.

For the Sphere a different method must be used so that all of the nodes are placed on the surface of the Sphere and the volume is vacant. For this geometry, I used the following equations:

$$x = \sqrt{1 - u^2} \cos \theta \quad (1)$$

$$y = \sqrt{1 - u^2} \sin \theta \quad (2)$$

$$z = u \quad (3)$$

where $\theta \in [0, 2\pi]$ and $u \in [-1, 1]$. This is guaranteed to uniformly distribute nodes on the surface area of the sphere [1].

All of these algorithms can be solved in $\Theta(n)$ where $n = |V|$ because each node only needs to be assigned a position once.

2.2.2 Edge Determination

There are several methods for finding the edges in the graph. The brute force method checks every node, and for each node checks all other nodes to see if they are close enough to form an edge. The brute force method is $\Theta(n^2)$ where $n = |V|$.

The second method to find the edges is the sweep method. This method first sorts the nodes along the x-axis. Then, for any node, we only need to search left and right until the distance along the x-axis is greater than the connection radius for the nodes. This dramatically reduces the search space. The sweep method is $O(n \lg(n) + 2rn^2)$ where $n = |V|$ and r is the connection radius. The $n \lg(n)$ portion is for the sorting and the $2rn^2$ portion is for measuring the distance between nodes in a sweep step.

The final method to find edges is the cell method. This method places the nodes into cells of area $r \times r$ based on their position in the topology. When the edge detection runs, each node needs to be visited once, but only the cell the node populates and the neighboring cells need to be searched for connections.

The only method that needs to be adjusted for the Sphere is the cell method. Instead of using a two dimensional grid of cells, a three dimensional mesh is needed to divide the topology. The cells then have volume $r \times r \times r$. Only the current cell and the neighboring cells need to be searched.

2.2.3 Graph Coloring

Two algorithms are used for coloring the graphs. The first is smallest-last vertex ordering, which sorts the vertices based on the number of degrees they have. The second is the greedy graph coloring algorithm.

Smallest-last vertex ordering is used to order the nodes for coloring. The steps to this algorithm are as follows [3]:

1. Initialize a representation of your target graph
2. Find the vertex v_j of minimum degree in your representation
3. Update your representation to simulate deleting v_j
4. If there are still vertices in the representation, return to step 1, otherwise terminate with the sequence of vertices removed

This algorithm is linear if each of the above steps is linear. Step 1 is linear if we can build a representation of the graph in linear time. For this, we can use an array of buckets, where each bucket holds the vertices that have the same number of edges as the position of the bucket in the array of buckets. To build this data structure, each node only needs to be visited once, making this linear in both space and time. Next, finding the vertex of minimum degree simply requires finding the lowest index bucket that has a node. This is bounded by the number of buckets, which is bounded by the number of nodes, making Step 2 linear. Next, we have to update the representation of the graph. To do this, we have to look at each node that shares an edge with v_j and move it to the bucket for nodes with one fewer degree. This requires traversing the list of edges for v_j which means Step 3 is linear. Since this is repeated for each node, the runtime of this program is $\Theta(|E| + |V|)$ and the space needed is $\Theta(|V|)$.

After this, a single traversal of the smallest-last vertex ordering is needed to color the graph. As we traverse this list, we check to see if the nodes before it (that are already colored) share an edge with the current node. The node can then be colored with any color it does not share an edge with or, if it shares an edge with all currently used colors, it is assigned a new color. This algorithm is also linear. Each node needs to be visited once and when a node is visited, all previous nodes are checked to see if they are in the edge list of the current node. Because we used smallest last vertex ordering, as we have to check more and more nodes, we get to check fewer and fewer edges. This makes the greedy coloring algorithm $O(|V| + |E|)$.

2.3 Algorithm Engineering

2.3.1 Node Placement

It is easy to implement the algorithms for placing nodes in the different geometries using Python's math library. This library offers functions for sampling points on a uniform distribution. For the Square, sampling on a range $[0, 1]$ is sufficient for all of the nodes. Since each node only needs to be placed once, this runs at $\Theta(n)$ where $n = |V|$.

For the Disk, the node needs to be resampled if it is too far from the center. To do this, the distance function is used to find the distance between the node and the center. If the node is further than 1 from the center, node generation falls into a while loop which iterates until the node is within the unit circle. Since nodes are taken from a uniform distribution, the number of nodes that will need to be resampled is approximately equal to the ratio of the area of the square that circumscribes the unit circle which falls outside of the unit circle to the total area of the square. This is given by:

$$\frac{(2r)^2 - \pi r^2}{(2r)^2} = \frac{4 - \pi}{4} = 0.2146 \quad (4)$$

Since the placement algorithm for each node of the Disk will iterate until the node falls within the unit circle, the total number of iterations N can be found as the sum of the geometric series:

$$N = \sum_{k=0}^{\infty} n(0.2146)^k = \frac{n}{1 - 0.2146} = 1.273n \quad (5)$$

where $n = |V|$. This shows this implementation is $\Theta(n)$.

For the node placement algorithm of the Sphere, again the math library in Python makes this easy. Each node needs two random values pulled from a uniform distribution, two square root operations, one sine operation, and one cosine operation. Each node only needs to be placed once so the runtime of this algorithm is $\Theta(n)$ where $n = |V|$.

2.3.2 Edge Determination

Each method implemented for finding edges has a different time complexity. The brute force method uses an outer loop and an inner loop, which each iterate over every node in the graph. An edge is saved to the adjacency list if the nodes are not the same and the distance between them is less than or equal to the calculated node radius. This is guaranteed to run in $\Theta(n^2)$ where $n = |V|$. The number of times the distance needs to be calculated is $n \times (n - 1)$ because it will not be calculated when the nodes are the same (distance would be zero, but no edge is drawn here). No additional space is needed for the brute force method so the space complexity is $O(1)$.

The implementation of sweep starts by sorting the nodes along the x-axis. Python lists have a built-in sort function that has $O(n \lg(n))$ time complexity [2]. After this stage, it iterates over every node building

a search space which will be scanned for edges. For each node, the list of nodes is searched left and right $r \times n$ nodes to find those within one radius length of the current node. With the search space built, the search space is iterated over once to find nodes that have a distance less than or equal the node radius. My implementation of this runs in $O(n \lg(n) + 4rn)$ where $n = |V|$ and r is the node connection radius. Because the list sort method sorts inplace, the only additional space needed is for the search space. This saves $O(2rn)$ nodes and is reset after every iteration.

The cell method implementation works in linear time. In the first step of the method, the cells are initialized as a list of empty lists. There are $(1/r + 1)^2$ cells. The nodes are then iterated over and assigned a cell by dividing their x and y coordinates by the node radius. At this point, the cells are iterated over and, for each node in the cell, the nodes in the current cell and the eight adjacent cells and the are checked to see if they fall within the node radius of the current node. All together, this implementation runs at $O(n + n + 9nr^2) = O((2 + 9r^2)n)$ where $n = |V|$. The amount of additional space needed is equal to the number of nodes because they are copied into their respective cells. This places the space complexity at $\Theta(n)$.

The cell method needs to be updated for the Sphere. To do this, an extra dimension is added to the cells, creating a 3D mesh. The only changes needed from the 2D method is that another loop is needed to iterate over the added dimension, and the search space turns into a 3x3 cube with the current cell at the center. Each node is still only visited once as the edges are determined. The runtime for this algorithm is $O(n + n + 27nr^3) = O((2 + 27r^3)n)$ where $n = |V|$. Again, the space complexity is $\Theta(n)$.

2.3.3 Graph Coloring

Implementing the smallest-last coloring algorithm involves implementing the smallest-last vertex ordering algorithm and the greedy graph coloring algorithm. For smallest-last vertex ordering, the first thing to do is build the data structure used to represent the graph with deleted nodes. The number of buckets needed it equal to the maximum degree of the nodes. Then, each node is placed in the bucket corresponding to the number of edges it has then the RGG. Simultaneously, a dictionary is created that maps each node to the number of degrees it has in the graph with deletions. Each value starts at the number of edges the corresponding node has in the RGG. At this point, we have iterated over all of the nodes once and allocated space for twice the number of nodes by copying them into the buckets and using them as the keys for the degrees dictionary.

Because Python dictionaries resize at specific numbers of entries, we can determine the number of additional insertions caused by rehashing while the degrees dictionary is built. Python dictionaries start out with space for 8 entries and quadruple in size until the number of entries is above 50,000, at which point it begins to double in size. Clearly the dictionary grows at a logarithmic rate, but the total number of insertions I for an input size of n is given by:

$$I = \begin{cases} n + 8 \sum_{k=1}^{\log_4 \lceil n/8 \rceil} 4^k & n \leq 50,000 \\ n + 8 \sum_{k=1}^6 4^k + 32768 \sum_{k=1}^{\log_2 \lceil n/32768 \rceil} 2^k & n > 50,000 \end{cases} \quad (6)$$

Fortunately, because the entire dictionary is built before it is used by the smallest-last vertex ordering algorithm, it will never again be resized once the algorithm starts. It will also be used to index into the buckets, so we gain a speed up here by not having to iterate over all of the edges for a node and determining if the node it shares an edge with are in the remaining graph each time we want to sift nodes down to lower buckets.

Next, the smallest-last vertex ordering algorithm is run until every node has been removed from the buckets. For each node, I iterate over the buckets from lowest degree to highest degree to find the first non-empty bucket. This bucket must contain the next node to remove because it contains all nodes with smallest degree. Before deleting the node from the graph and moving all adjacent nodes down a bucket, I check to see if the current bucket has all remaining nodes. If this is the case, the terminal clique has been found, and the size of the terminal clique must be saved. After this check, a node is popped from the end of the current bucket, and appended to the smallest-last ordering result. Then, for all the adjacent nodes to the popped node in the original graph, I try to remove that node from the bucket with its degree. If the delete fails, then the node has already been removed. Otherwise, the number of degrees for that node can be decremented and the node can be appended to the correct bucket for its new degree.

The last step is to reverse the order of the smallest-last ordering result because it was built in the opposite order (smallest-first). All together, excluding the initialization of accessory data structures, this implementation runs in $\Theta(|V| + |V| \times |E|)$ time and $\Theta(2|V|)$ space since nodes are removed from the buckets and added to the result.

After this the graph needs to be colored. For this I initially assign each node a color of 0. I iterate over all of the nodes in the smallest-last vertex ordering. At each node, I generate a list of colors that is already used by the neighbors of that node by iterating over all of the nodes before it in the smallest-last ordering and checking if they exist in the list of edges for the current node. Then, I just have to increment color from 0 until it does not exist in the search space and I have the color to assign to the node.

Since the smallest-last ordering is used, each time I check to see if a node is adjacent to the current node, I am searching nodes with fewer and fewer edges. This means that the nodes with the most neighbors are searched first, when the number of other nodes to check is lowest, and the nodes with the fewest neighbors are searched last, when we have the most nodes to check if they share an edge with the current node. All together, this implementation runs in $\Theta(1.5|V| + |E|)$ time and $\Theta(|V|)$ space because we need a new array for the colors assigned to each of the nodes.

2.4 Verification

2.4.1 Node Placement

The nodes can be verified to be distributed uniformly if the distribution of degrees follows a normal distribution. To show that the distribution of degrees for each of my geometries are following a normal distribution, I plotted degree histograms for each of the geometries with 32,000 nodes and an average degree of 16. The histogram for Square is given in Figure 1, Disk is given in Figure 2, and Sphere is given in Figure 3. These histograms clearly follow a normal distribution.

2.4.2 Edge Determination

The runtime for the edge detection methods can be verified by varying the number of nodes and measuring the runtime of each algorithm. By looking at how the runtime grows, we can calculate the trendline that best fits the growth rate. For the first comparison, I vary the number of nodes from 4,000 to 64,000 in steps of 4,000, while holding the desired average degree constant at 16. As we can see in Figure 4, the growth rates of the brute force and sweep methods are quadratic, while the growth rate of the cell method. The trendline functions are given on the graph.

For the second metric, I held the number of nodes constant at 32,000 and varied the desired average degree from 2 to 32 in steps of 2. The graph is given in Figure 5. The cell method clearly grows linearly, but the sweep method is harder to gauge. Since varying the desired average degree should only change the node radius, I would expect this to grow linearly as well. However, because each graph is randomly generated, some graphs can have nodes that are closer to sorted order than others. This can effect the measured runtime. It would be easier to gauge the trend if it I ran the data collection multiple times and averaged the results.

References

- [1] Weisstein, Eric W., Wolfram MathWorld
Sphere Point Picking
<http://mathworld.wolfram.com/SpherePointPicking.html>
- [2] Peters, Tim
Timsort
<http://svn.python.org/projects/python/trunk/Objects/listsort.txt>
- [3] Matula, David and Beck, Leland
Smallest-Last Ordering and Clustering and Graph Coloring Algorithms
- [4] Johnson, Ian
Linear-Time Computation of High-Converage Backbones for Wireless Sensor Networks
<https://github.com/ianjohnson/SensorNetwork/blob/master/Report/Report.pdf>

3 Appendix A - Figures

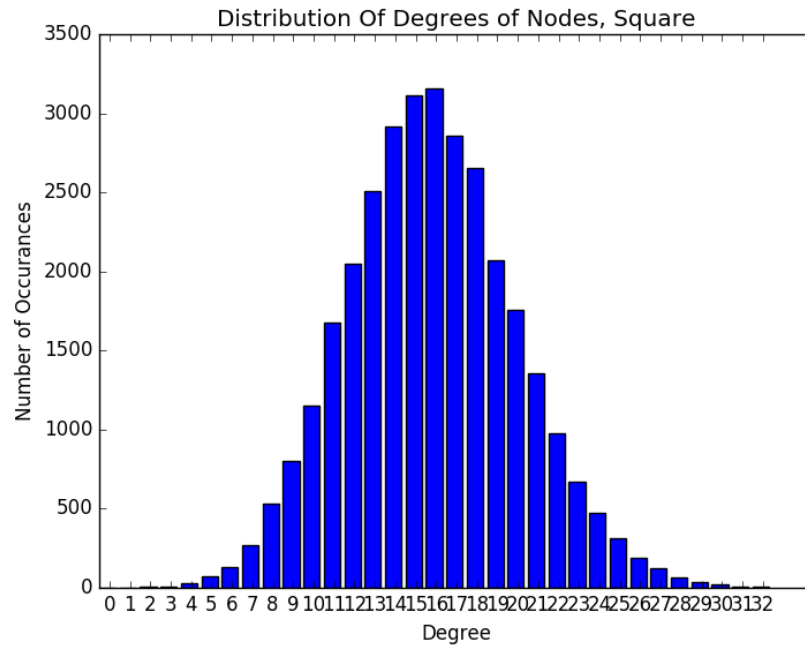


Figure 1: Distribution of Degree counts for Square. 32,000 Nodes, Average Degree of 16

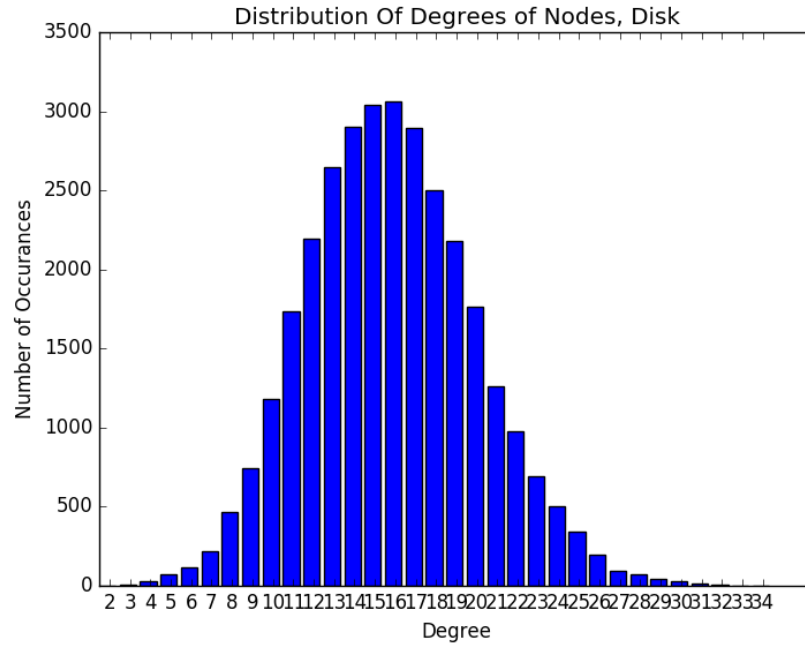


Figure 2: Distribution of Degree counts for Disk. 32,000 Nodes, Average Degree of 16

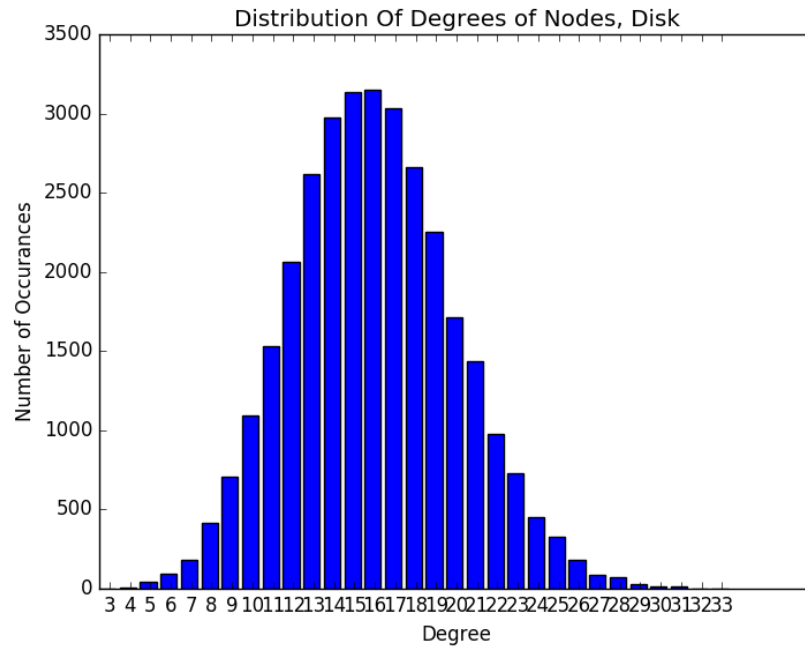


Figure 3: Distribution of Degree counts for Sphere. 32,000 Nodes, Average Degree of 16

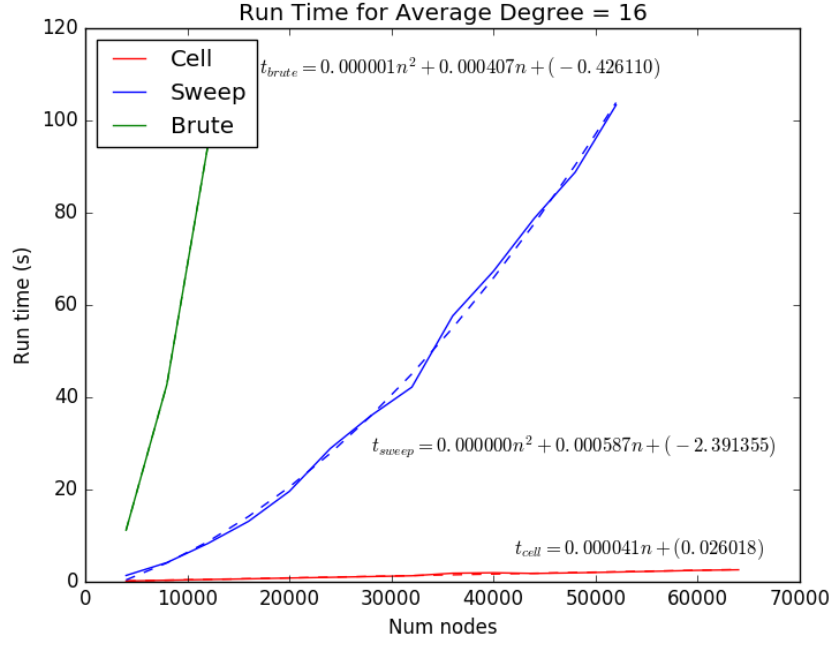


Figure 4: Runtime for Each Edge Detection Method, Average Degree of 16

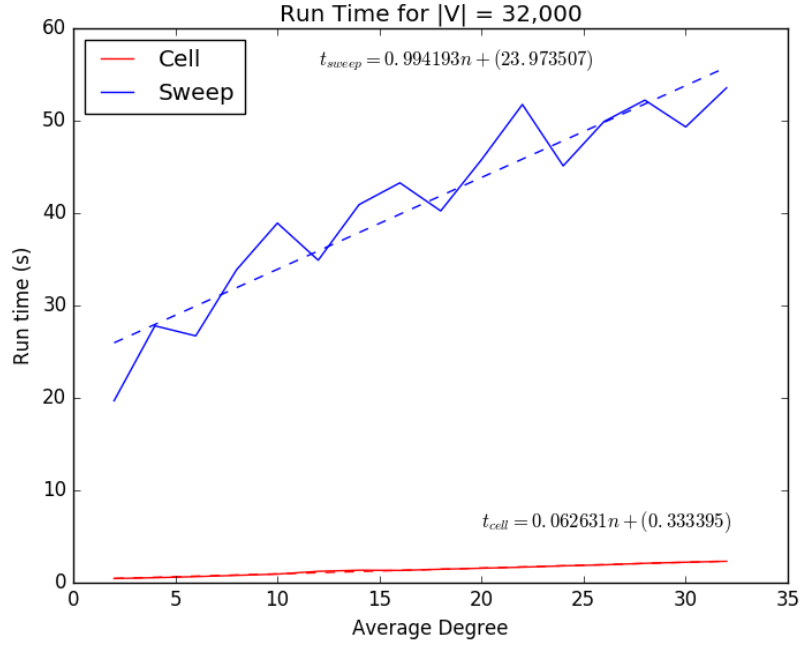


Figure 5: Runtime for Cell and Sweep Edge Detection, Variable Average Degree

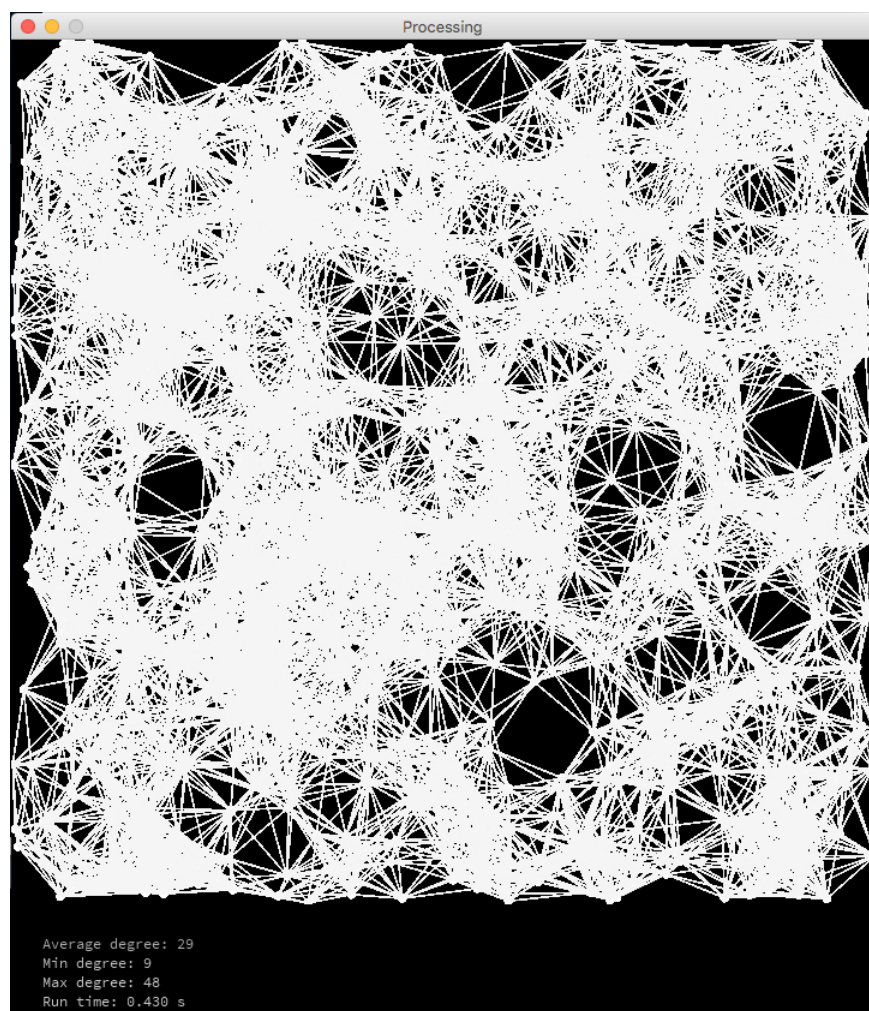


Figure 6: Square Benchmark Number 1. 1000 Nodes, Average Degree of 32

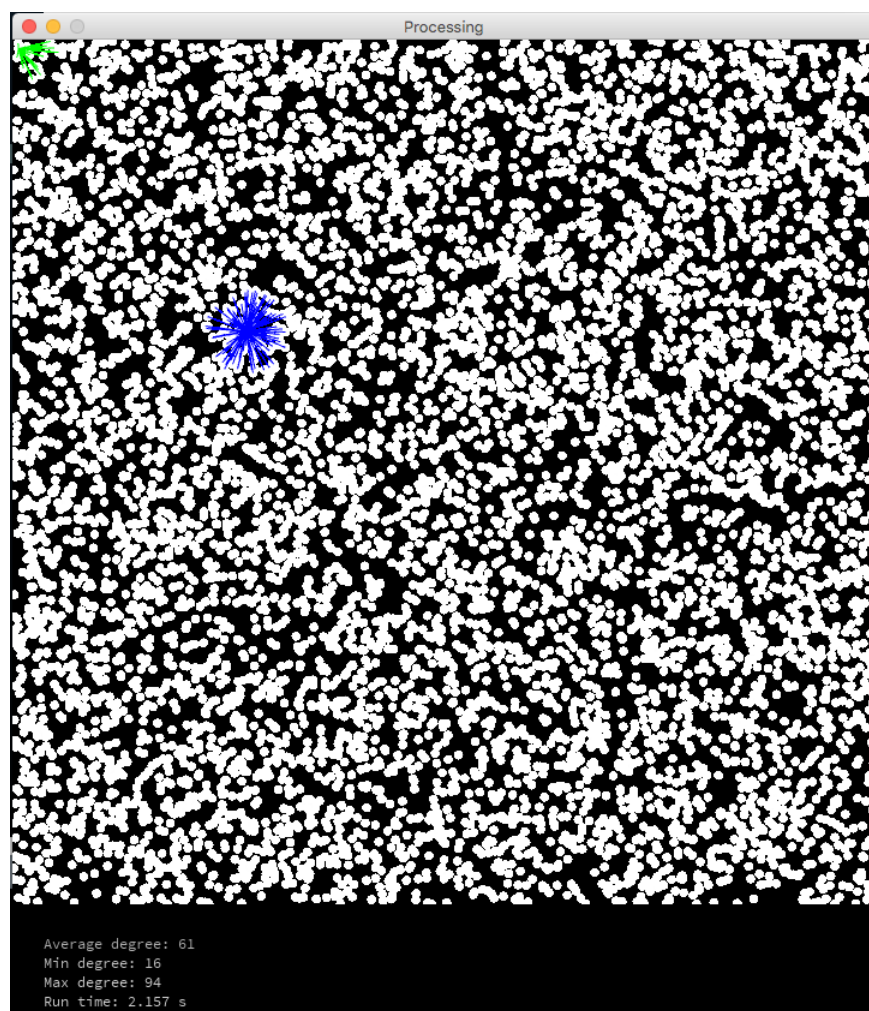


Figure 7: Square Benchmark Number 2. 8000 Nodes, Average Degree of 64

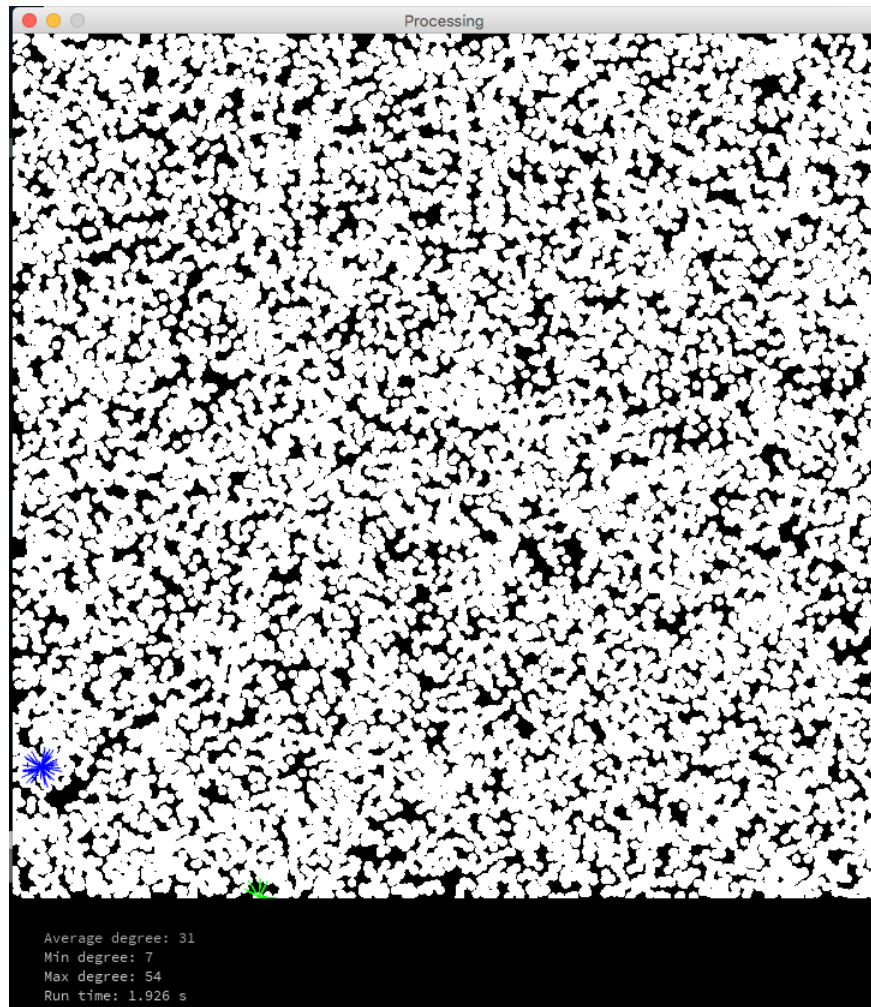


Figure 8: Square Benchmark Number 3. 16000 Nodes, Average Degree of 32




Figure 9: Square Benchmark Number 4. 64000 Nodes, Average Degree of 64

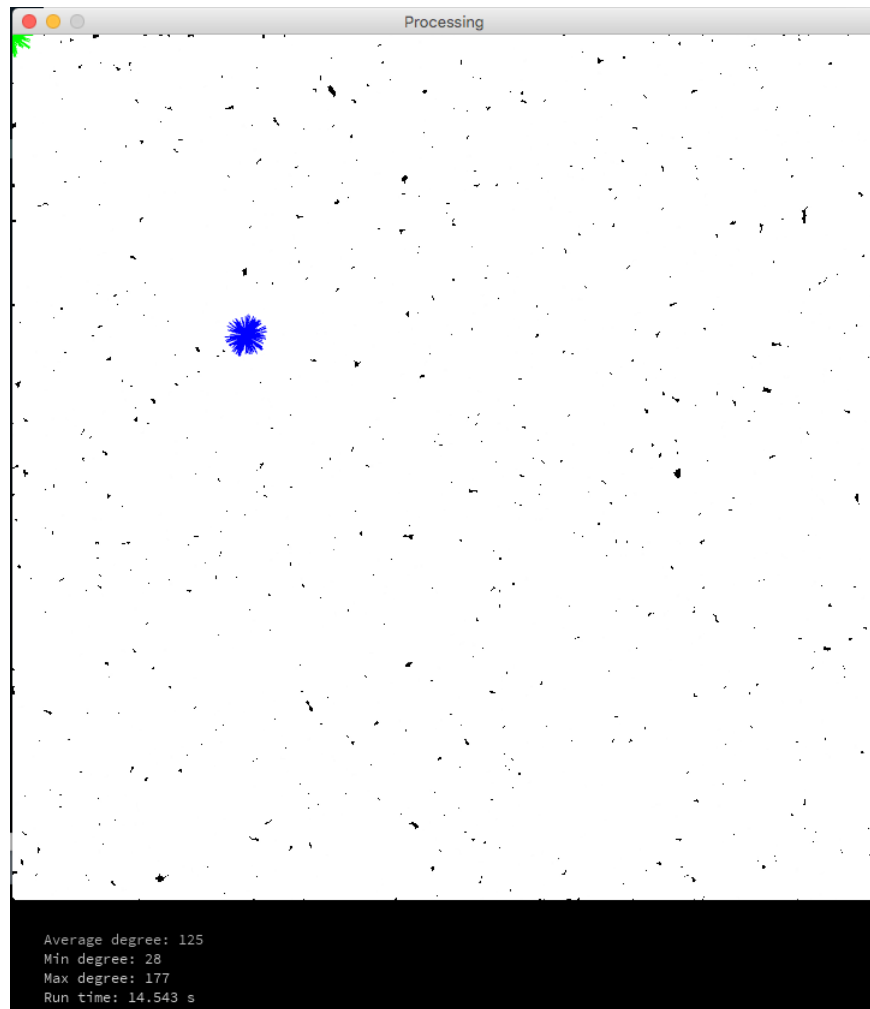


Figure 10: Square Benchmark Number 5. 64000 Nodes, Average Degree of 128

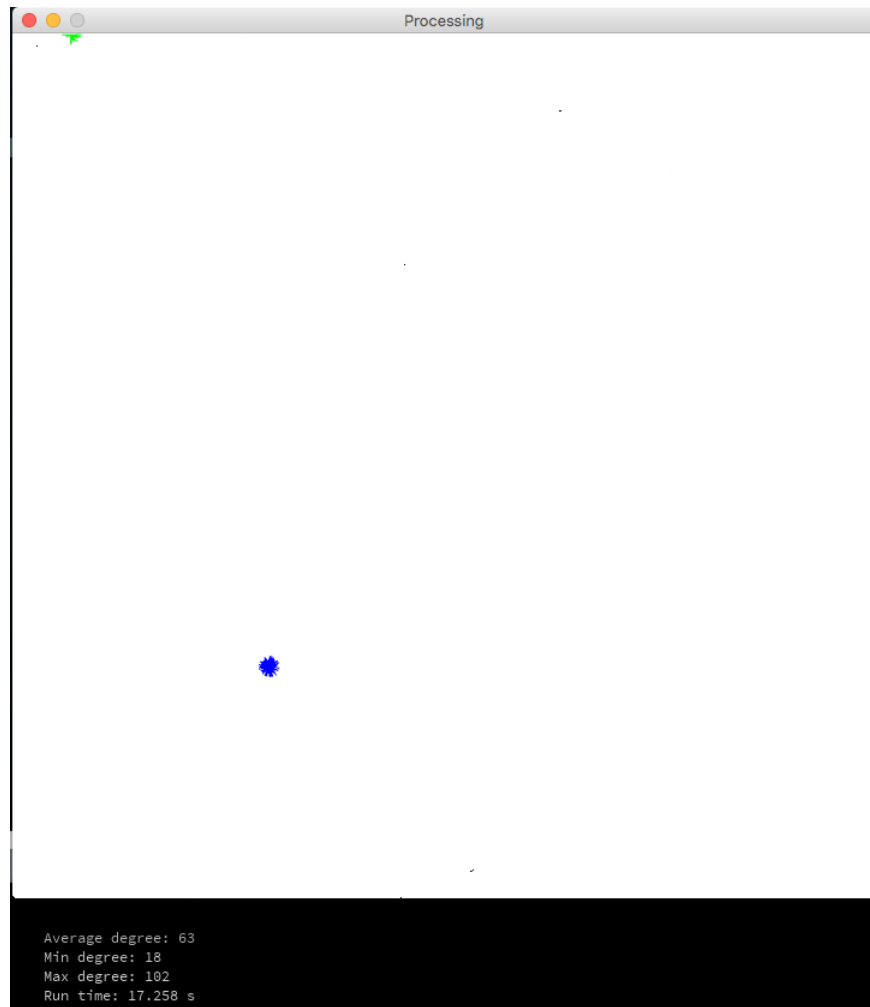


Figure 11: Square Benchmark Number 6. 128000 Nodes, Average Degree of 64

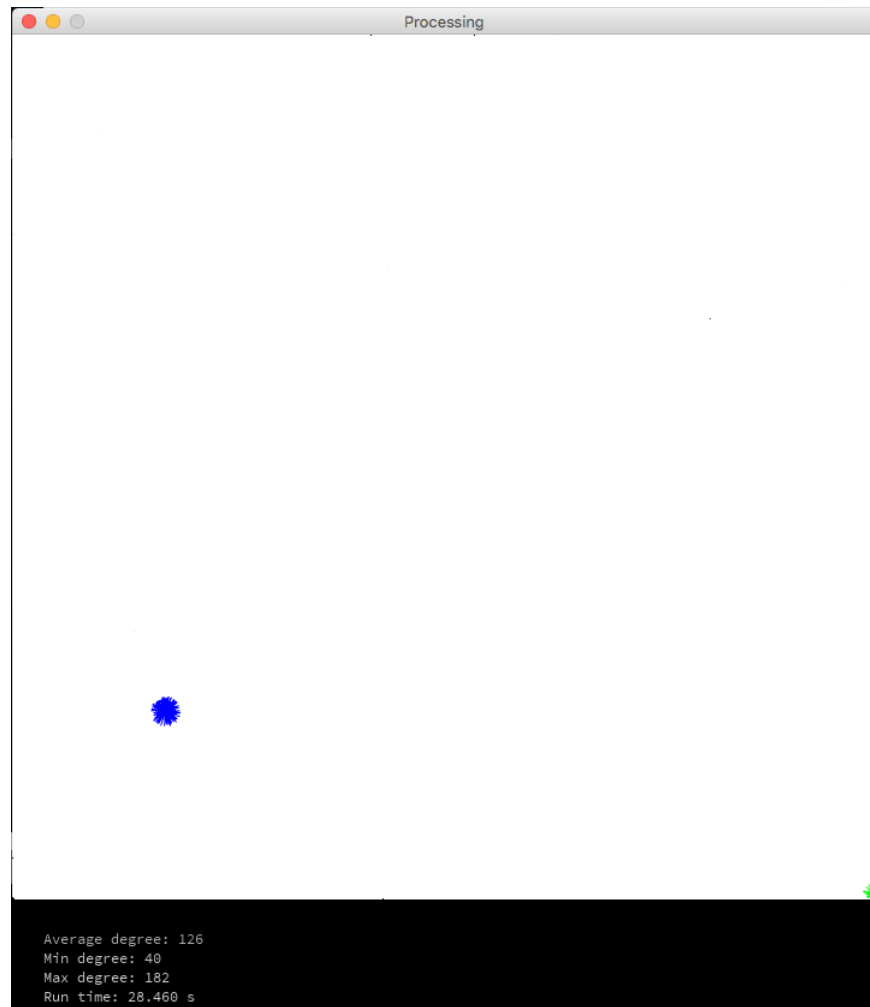


Figure 12: Square Benchmark Number 7. 128000 Nodes, Average Degree of 128

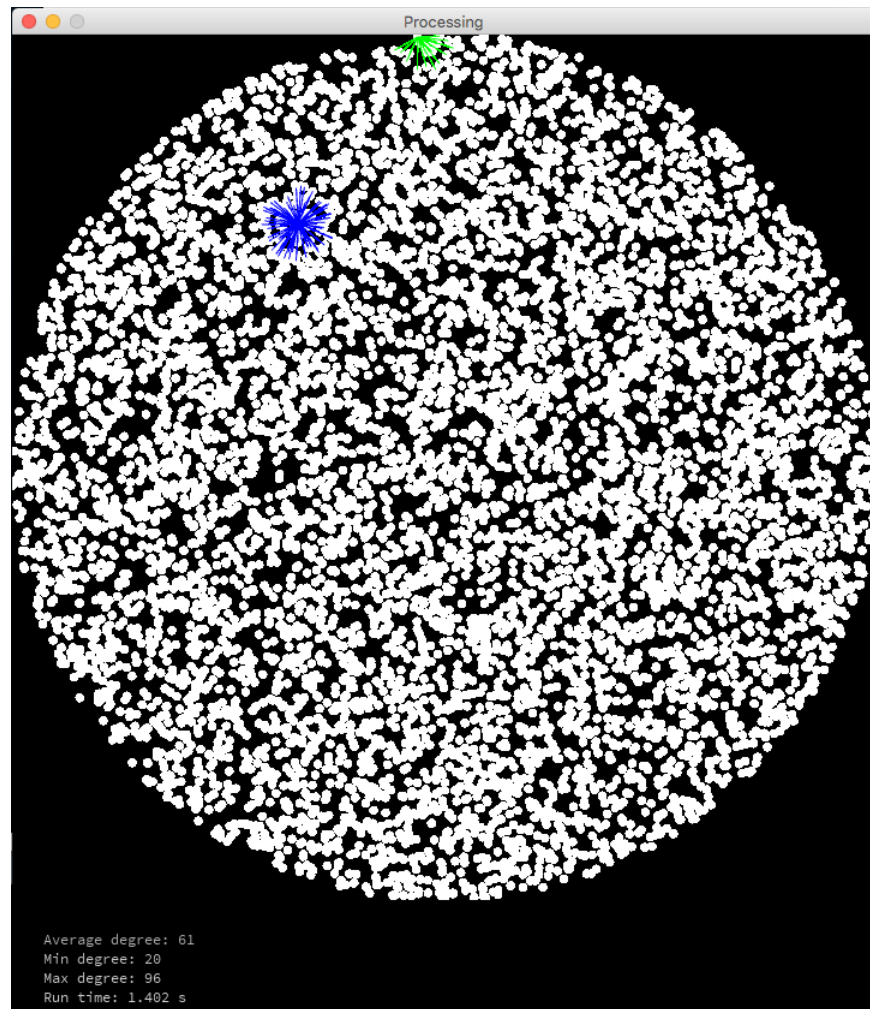


Figure 13: Disk Benchmark Number 1. 8000 Nodes, Average Degree of 64

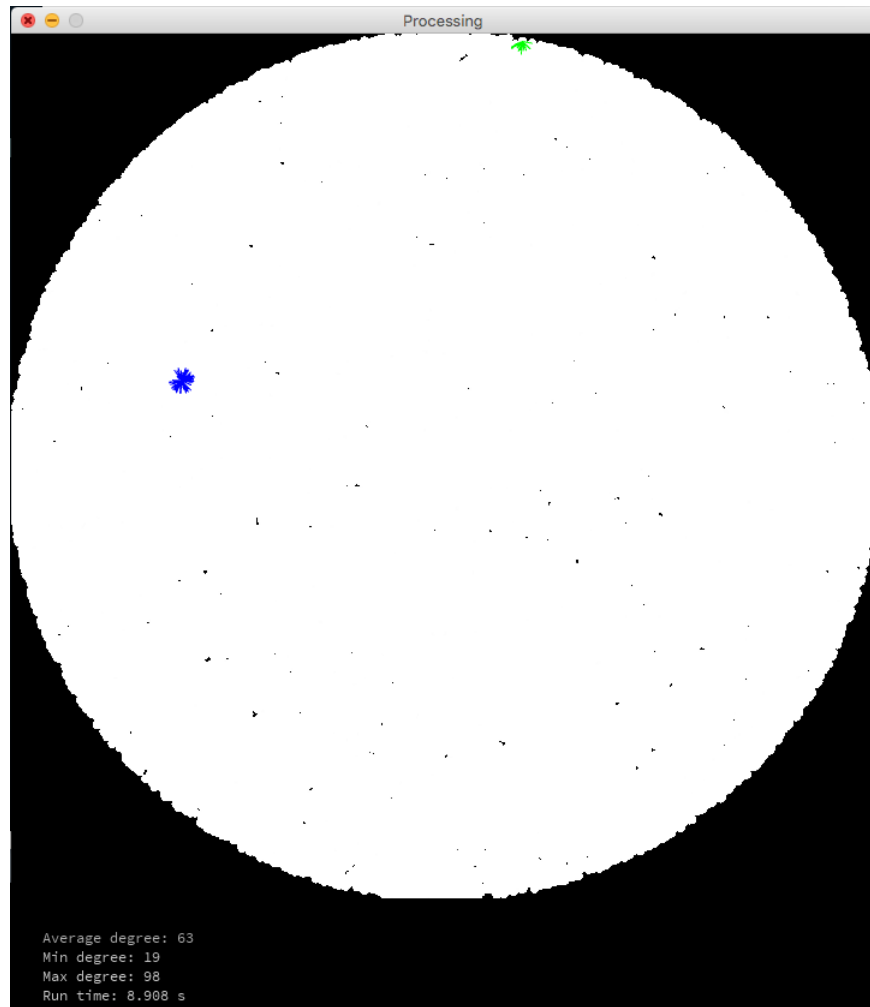


Figure 14: Disk Benchmark Number 2. 64000 Nodes, Average Degree of 64

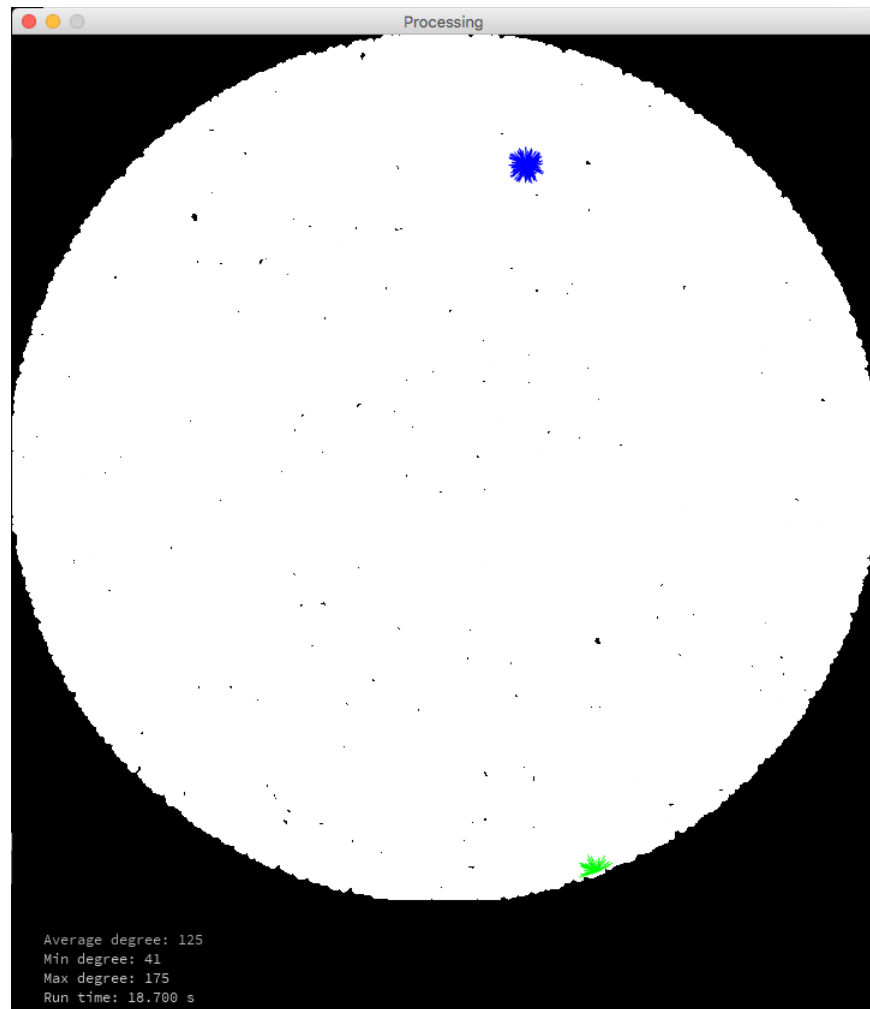


Figure 15: Disk Benchmark Number 3. 64000 Nodes, Average Degree of 128

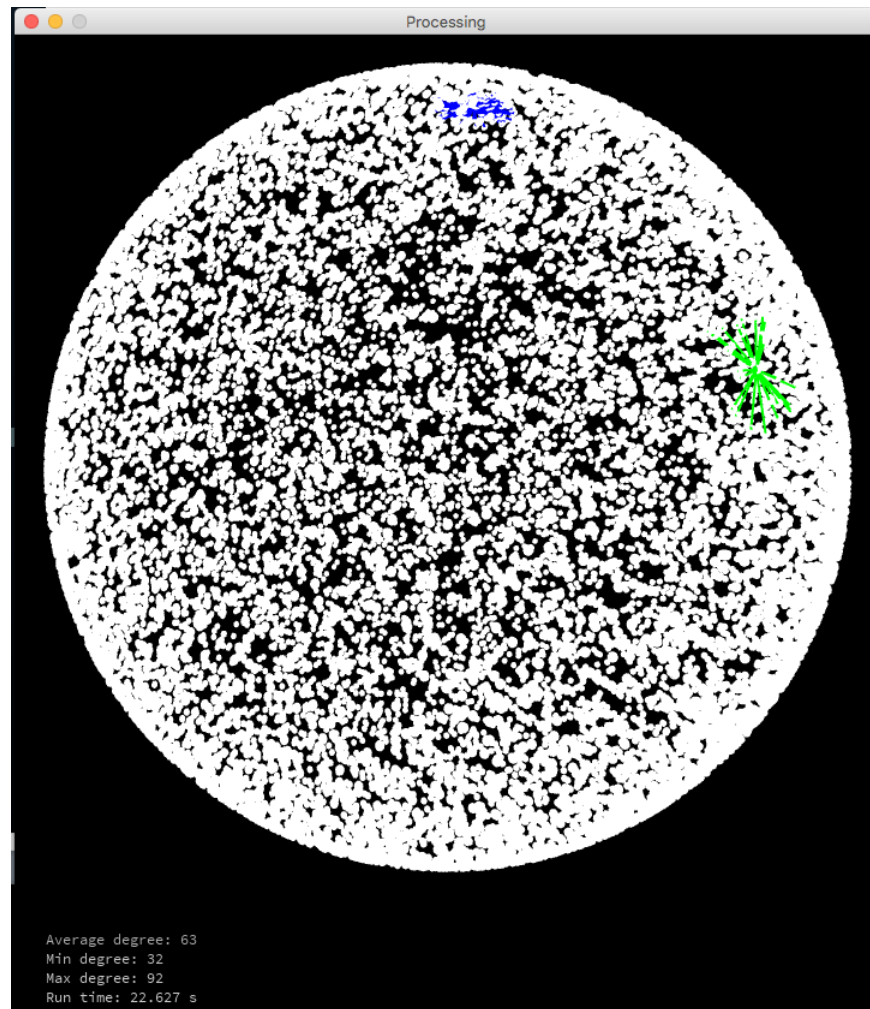


Figure 16: Sphere Benchmark Number 1. 16000 Nodes, Average Degree of 64

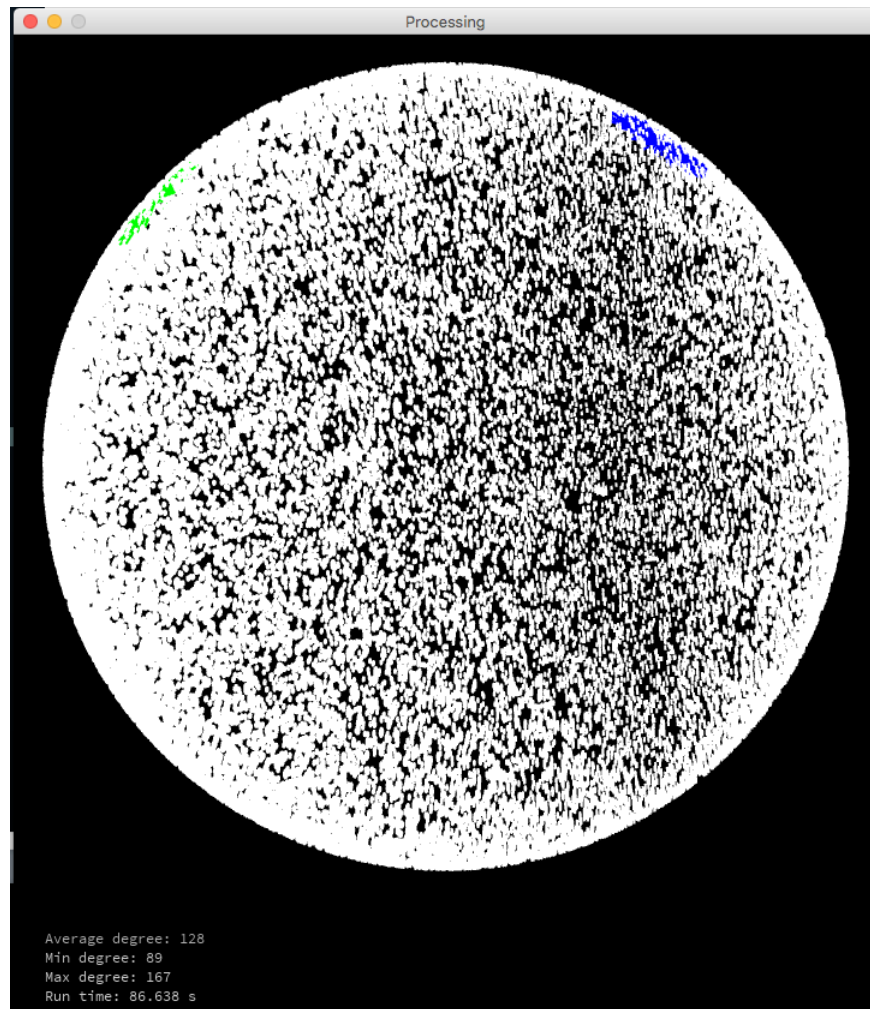


Figure 17: Sphere Benchmark Number 2. 32000 Nodes, Average Degree of 128

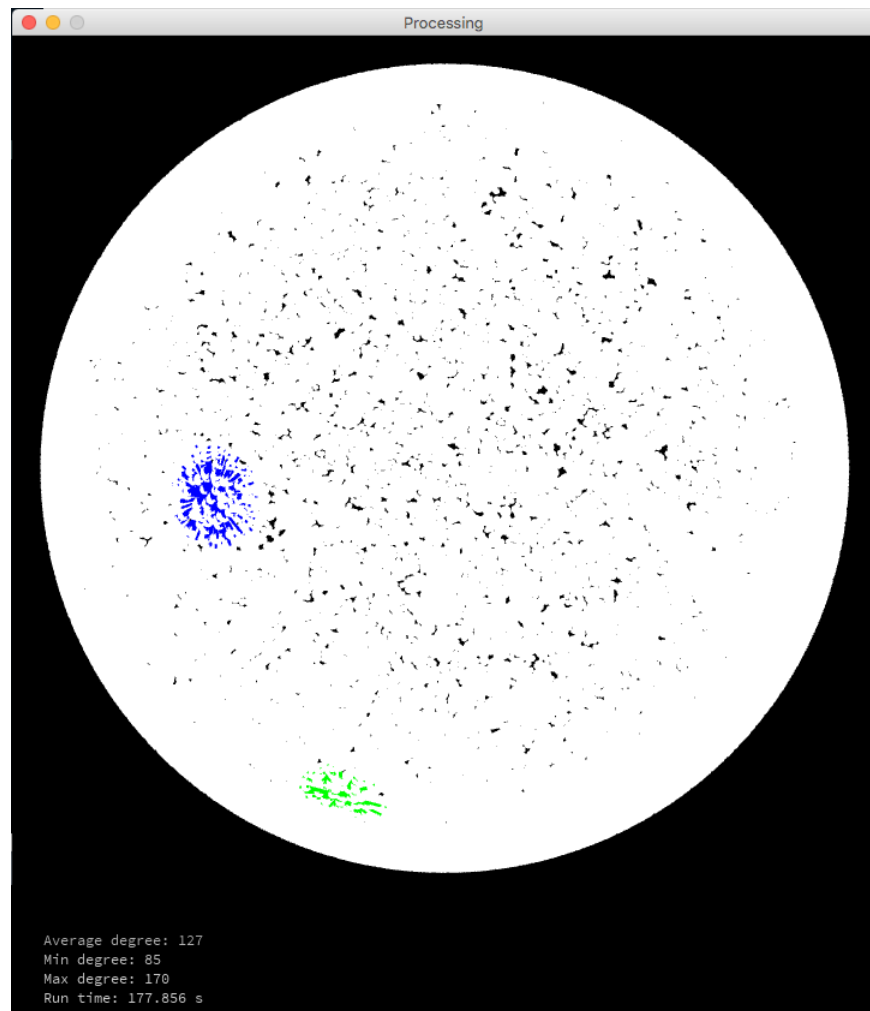


Figure 18: Sphere Benchmark Number 3. 64000 Nodes, Average Degree of 128

4 Appendix B - Code Listings

Listing 1: Processing driver

```
1 import random
2 import time
3 import math
4 from collections import Counter
5 from objects.topology import Square, Disk, Sphere
6
7 CANVAS_HEIGHT = 720
8 CANVAS_WIDTH = 720
9
10 NUMNODES = 1000
11 AVG_DEG = 32
12
13 MAX_NODES_TO_DRAW_EDGES = 8000
14
15 RUN_BENCHMARK = False
16
17 def setup():
18     size(CANVAS_WIDTH, CANVAS_HEIGHT, P3D)
19     background(0)
20
21 def draw():
22     topology.drawGraph(MAX_NODES_TO_DRAW_EDGES)
23
24 def main():
25     global topology
26     # topology = Square()
27     # topology = Disk()
28     topology = Sphere()
29
30     topology.num_nodes = NUMNODES
31     topology.avg_deg = AVG_DEG
32     topology.canvas_height = CANVAS_HEIGHT
33     topology.canvas_width = CANVAS_WIDTH
34
35     if RUN_BENCHMARK:
36         n_benchmark = 0
37         topology.prepBenchmark(n_benchmark)
38
39     run_time = time.clock()
40
41     topology.generateNodes()
42     topology.findEdges(method="cell")
43     topology.colorGraph()
44
45     print "Average degree: {}".format(topology.findAvgDegree())
46     print "Min degree: {}".format(topology.getMinDegree())
47     print "Max degree: {}".format(topology.getMaxDegree())
48     print "Terminal clique size: {}".format(topology.term_clique_size)
49     print "Number of colors: {}".format(len(set(topology.node_colors)))
50     print "Max degree when deleted: {}".format(max(topology.deg_when_del.values()))
51
52     color_cnt = Counter(topology.node_colors)
53     print "Max color set size: {} color: {}".format(color_cnt.most_common(1)
54     [0][1],
55     color_cnt.most_common(1)
56     [0][0])
57
58     run_time = time.clock() - run_time
59     print "Run time: {:.3f} s".format(run_time)
60
61 main()
```

Listing 2: Topology class and subclasses

```

1 import random
2 import math
3 import time
4
5 # benchmarks (num_nodes, avg_deg)
6 SQUARE_BENCHMARKS = [(1000,32), (8000,64), (16000,32), (64000,64), (64000,128),
7                       (128000,64), (128000, 128)]
8 DISK_BENCHMARKS = [(8000,64), (64000,64), (64000,128)]
9 SPHERE_BENCHMARKS = [(16000,64), (32000,128), (64000,128)]
10
11 """
12 Topology – super class for the shape of the random geometric graph
13 """
14 class Topology(object):
15
16     num_nodes = 100
17     avg_deg = 0
18     canvas_height = 720
19     canvas_width = 720
20
21     def __init__(self):
22         self.nodes = []
23         self.edges = {}
24         self.node_r = 0.0
25         self.minDeg = ()
26         self.maxDeg = ()
27         self.s_last = []
28         self.deg_when_del = {}
29         self.node_colors = []
30
31     # public function for generating nodes of the graph, must be subclassed
32     def generateNodes(self):
33         print "Method for generating nodes not subclassed"
34
35     # public function for finding edges
36     def findEdges(self, method="brute"):
37         self._getRadiusForAverageDegree()
38         self._addNodesAsEdgeKeys()
39
40         if method == "brute":
41             self._bruteForceFindEdges()
42         elif method == "sweep":
43             self._sweepFindEdges()
44         elif method == "cell":
45             self._cellFindEdges()
46         else:
47             print "Find edges method not defined: {}".format(method)
48
49         self._findMinAndMaxDegree()
50
51     # brute force edge detection
52     def _bruteForceFindEdges(self):
53         for i, n in enumerate(self.nodes):
54             for j, m in enumerate(self.nodes):
55                 if i != j and self._distance(n, m) <= self.node_r:
56                     self.edges[n].append(j)
57
58     # sweep edge detection (2D)
59     def _sweepFindEdges(self):
60         # TODO: Only look forward
61         self.nodes.sort(key=lambda x: x[0])
62
63         for i, n in enumerate(self.nodes):
64             search_space = []
65             for j in range(1, i+1):
66                 if abs(n[0] - self.nodes[i-j][0]) <= self.node_r:
67                     search_space.append(i-j)

```

```

68         else:
69             break
70     for j in range(1, self.num_nodes-i):
71         if abs(n[0] - self.nodes[i+j][0]) <= self.node_r:
72             search_space.append(i+j)
73         else:
74             break
75     for j in search_space:
76         if self._distance(n, self.nodes[j]) <= self.node_r:
77             self.edges[n].append(j)
78
79 # cell edge detection (2D)
80 def _cellFindEdges(self):
81     num_cells = int(1/self.node_r) + 1
82     cells = []
83     for i in range(num_cells):
84         cells.append([[j for j in range(num_cells)]])
85
86     for i, n in enumerate(self.nodes):
87         cells[int(n[0]/self.node_r)][int(n[1]/self.node_r)].append(i)
88
89     for i in range(num_cells):
90         for j in range(num_cells):
91             for n_i in cells[i][j]:
92                 for c in self._findAdjCells(i, j, num_cells):
93                     for m_i in cells[c[0]][c[1]]:
94                         if self._distance(self.nodes[n_i], self.nodes[m_i]) <=
self.node_r and n_i != m_i:
95                             self.edges[self.nodes[n_i]].append(m_i)
96
97 # cell edge detection helper function (2D)
98 def _findAdjCells(self, i, j, n):
99     # TODO: Only look forward
100     xRange = [(i-1)%n, i, (i+1)%n]
101     yRange = [(j-1)%n, j, (j+1)%n]
102     return ((x,y) for x in xRange for y in yRange)
103
104 # function for finding the radius needed for the desired average degree
105 # must be subclassed
106 def _getRadiusForAverageDegree(self):
107     print "Method for finding necessary radius for average degree not
subclassed"
108
109 # helper function for findEdges, initializes edges dict
110 def _addNodesAsEdgeKeys(self):
111     self.edges = {n:[] for n in self.nodes}
112
113 # calculates the distance between two nodes (2D)
114 def _distance(self, n, m):
115     return math.sqrt((n[0] - m[0])**2+(n[1] - m[1])**2)
116
117 # public function for finding the number of edges
118 def findNumEdges(self):
119     sigma_edges = 0
120     for k in self.edges.keys():
121         sigma_edges += len(self.edges[k])
122
123     return sigma_edges/2
124
125 # public function for finding the average degree of nodes
126 def findAvgDegree(self):
127     return 2*self.findNumEdges()/self.num_nodes
128
129 # helper function for finding nodes with min and max degree
130 def _findMinAndMaxDegree(self):
131     self.minDeg = self.edges.keys()[0]
132     self.maxDeg = self.edges.keys()[0]
133

```

```

134         for k in self.edges.keys():
135             if len(self.edges[k]) < len(self.edges[self.minDeg]):
136                 self.minDeg = k
137             if len(self.edges[k]) > len(self.edges[self.maxDeg]):
138                 self.maxDeg = k
139
140     # public function for getting the minimum degree
141     def getMinDegree(self):
142         return len(self.edges[self.minDeg])
143
144     # public functino for getting the maximum degree
145     def getMaxDegree(self):
146         return len(self.edges[self.maxDeg])
147
148     # public function for setting up the benchmark to run, must be subclassed
149     def prepBenchmark(self, n):
150         print "Method for preparing benchmark not subclassed"
151
152     # public function for drawing the graph
153     def drawGraph(self, n_limit):
154         self._drawNodes()
155         if self.num_nodes <= n_limit:
156             self._drawEdges()
157         else:
158             self._drawMinMaxDegNodes()
159
160     # responsible for drawing the nodes in the canvas
161     def _drawNodes(self):
162         strokeWeight(2)
163         stroke(255)
164         fill(255)
165
166         for n in range(self.num_nodes):
167             ellipse(self.nodes[n][0]*self.canvas_width, self.nodes[n][1]*self.
canvas_height, 5, 5)
168
169     # responsible for drawing the edges in the canvas
170     def _drawEdges(self):
171         strokeWeight(1)
172         stroke(245)
173         fill(255)
174
175         for n in self.edges.keys():
176             for m_i in self.edges[n]:
177                 line(n[0]*self.canvas_width, n[1]*self.canvas_height, self.nodes[
m_i][0]*self.canvas_width, self.nodes[m_i][1]*self.canvas_height)
178
179     # responsible for drawing the edges of the min and max degree nodes
180     def _drawMinMaxDegNodes(self):
181         strokeWeight(1)
182         stroke(0,255,0)
183         fill(255)
184         for n_i in self.edges[self.minDeg]:
185             line(self.minDeg[0]*self.canvas_width, self.minDeg[1]*self.
canvas_height, self.nodes[n_i][0]*self.canvas_width, self.nodes[n_i][1]*self.
canvas_height)
186
187             stroke(0,0,255)
188             for n_i in self.edges[self.maxDeg]:
189                 line(self.maxDeg[0]*self.canvas_width, self.maxDeg[1]*self.
canvas_height, self.nodes[n_i][0]*self.canvas_width, self.nodes[n_i][1]*self.
canvas_height)
190
191     # uses smallest last vertex ordering to color the graph
192     def colorGraph(self):
193         self.s_last, self.deg_when_del = self._smallestLastVertexOrdering()
194         self.node_colors = self._assignNodeColors(self.s_last)
195

```

```

196 # constructs a degree structure and determines the smallest last vertex
197 ordering
198 def _smallestLastVertexOrdering(self):
199     deg_sets = {l:set() for l in range(len(self.edges[self.maxDeg])+1)}
200     deg_when_del = {n:len(self.edges[n]) for n in self.nodes}
201
202     for i, n in enumerate(self.nodes):
203         deg_sets[deg_when_del[n]].add(i)
204
205     smallest_last_ordering = []
206
207     clique_found = False
208     j = len(self.nodes)
209     while j > 0:
210         # get the current smallest bucket
211         curr_bucket = 0
212         while len(deg_sets[curr_bucket]) == 0:
213             curr_bucket += 1
214
215         # if all the remaining nodes are connected we have the terminal clique
216         if not clique_found and len(deg_sets[curr_bucket]) == j:
217             clique_found = True
218             self.term_clique_size = curr_bucket
219
220         # get node with smallest degree
221         v_i = deg_sets[curr_bucket].pop()
222         smallest_last_ordering.append(v_i)
223
224         # decrement position of nodes that shared an edge with v
225         for n_i in (n_i for n_i in self.edges[self.nodes[v_i]] if n_i in
226 deg_sets[deg_when_del[self.nodes[n_i]]]):
227             deg_sets[deg_when_del[self.nodes[n_i]]].remove(n_i)
228             deg_when_del[self.nodes[n_i]] -= 1
229             deg_sets[deg_when_del[self.nodes[n_i]]].add(n_i)
230
231         j -= 1
232
233     # reverse list since it was built shortest-first
234     return smallest_last_ordering[::-1], deg_when_del
235
236 # assigns the colors to nodes given in a smallest-last vertex ordering as a
237 parallel array
238 def _assignNodeColors(self, s_last):
239     colors = [-1 for _ in range(len(s_last))]
240     for i in s_last:
241         adj_colors = set([colors[j] for j in self.edges[self.nodes[i]]])
242         color = 0
243         while color in adj_colors:
244             color += 1
245         colors[i] = color
246
247     return colors
248
249 """
250 Square — inherits from Topology, overloads generateNodes and
251 _getRadiusForAverageDegree
252 for a unit square topology
253 """
254 class Square(Topology):
255
256     def __init__(self):
257         super(Square, self).__init__()
258
259     # places nodes uniformly in a unit square
260     def generateNodes(self):
261         for i in range(self.num_nodes):
262             self.nodes.append((random.uniform(0,1), random.uniform(0,1)))

```

```

260     # calculates the radius needed for the requested average degree in a unit
261     # square
262     def _getRadiusForAverageDegree(self):
263         self.node_r = math.sqrt(self.avg_deg/(self.num_nodes * math.pi))
264
265     # gets benchmark setting for square
266     def prepBenchmark(self, n):
267         self.num_nodes = SQUARE_BENCHMARKS[n][0]
268         self.avg_deg = SQUARE_BENCHMARKS[n][1]
269
270     """
271     Disk – inherits from Topology, overloads generateNodes and
272     _getRadiusForAverageDegree
273     for a unit circle topology
274     """
275     class Disk(Topology):
276
277         def __init__(self):
278             super(Disk, self).__init__()
279
280         # places nodes uniformly in a unit square and regenerates the node if it falls
281         # outside of the circle
282         def generateNodes(self):
283             for i in range(self.num_nodes):
284                 p = (random.uniform(0,1), random.uniform(0,1))
285                 while self._distance(p, (0.5,0.5)) > 0.5:
286                     p = (random.uniform(0,1), random.uniform(0,1))
287                 self.nodes.append(p)
288
289         # calculates the radius needed for the requested average degree in a unit
290         # circle
291         def _getRadiusForAverageDegree(self):
292             self.node_r = math.sqrt((self.avg_deg + 0.0)/self.num_nodes)/2
293
294         # gets benchmark setting for disk
295         def prepBenchmark(self, n):
296             self.num_nodes = DISK_BENCHMARKS[n][0]
297             self.avg_deg = DISK_BENCHMARKS[n][1]
298
299     """
300     Sphere – inherits from Topology, overloads generateNodes,
301     _getRadiusForAverageDegree,
302     and _distance for a unit sphere topology. Also updates the drawGraph function for
303     a 3D canvas
304     """
305     class Sphere(Topology):
306
307         # adds rotation and node limit variables
308         def __init__(self):
309             super(Sphere, self).__init__()
310             self.rot = (0,math.pi/4,0) # this may move to Topology if rotation is
311             # given to the 2D shapes
312             # used to control _drawNodes functionality
313             self.n_limit = 8000
314
315         # places nodes in a unit cube and projects them onto the surface of the sphere
316         def generateNodes(self):
317             for i in range(self.num_nodes):
318                 # equations for uniformly distributing nodes on the surface area of
319                 # a sphere: http://mathworld.wolfram.com/SpherePointPicking.html
320                 u = random.uniform(-1,1)
321                 theta = random.uniform(0, 2*math.pi)
322                 p = (
323                     math.sqrt(1 - u**2) * math.cos(theta),
324                     math.sqrt(1 - u**2) * math.sin(theta),
325                     u
326                 )
327                 self.nodes.append(p)

```

```

323
324 # overrides cell for 3D topology, uses 3D mesh of buckets
325 def _cellFindEdges(self):
326     num_cells = int(1/self.node_r) + 1
327     cells = []
328     for i in range(num_cells):
329         cells.append([[] for k in range(num_cells)] for j in range(num_cells)
330 ])
331
332     for i, n in enumerate(self.nodes):
333         cells[int(n[0]/self.node_r)][int(n[1]/self.node_r)][int(n[2]/self.
334 node_r)].append(i)
335
336     for i in range(num_cells):
337         for j in range(num_cells):
338             for k in range(num_cells):
339                 for n_i in cells[i][j][k]:
340                     for c in self._findAdjCells(i, j, k, num_cells):
341                         for m_i in cells[c[0]][c[1]][c[2]]:
342                             if self._distance(self.nodes[n_i], self.nodes[m_i
343 ]) <= self.node_r and n_i != m_i:
344                                 self.edges[self.nodes[n_i]].append(m_i)
345
346 # overrides adjacent cell finding for 3x3 surrounding buckets
347 def _findAdjCells(self, i, j, k, n):
348     # TODO: Only look forward
349     xRange = [(i-1)%n, i, (i+1)%n]
350     yRange = [(j-1)%n, j, (j+1)%n]
351     zRange = [(k-1)%n, k, (k+1)%n]
352     return ((x,y,z) for x in xRange for y in yRange for z in zRange)
353
354 # calculates the radius needed for the requested average degree in a unit
355 sphere
356 def _getRadiusForAverageDegree(self):
357     self.node_r = math.sqrt((self.avg_deg + 0.0)/self.num_nodes)*2
358
359 # calculates the distance between two nodes (3D)
360 def _distance(self, n, m):
361     return math.sqrt((n[0] - m[0])**2+(n[1] - m[1])**2+(n[2] - m[2])**2)
362
363 # gets benchmark setting for sphere
364 def prepBenchmark(self, n):
365     self.num_nodes = SPHERE_BENCHMARKS[n][0]
366     self.avg_deg = SPHERE_BENCHMARKS[n][1]
367
368 # public function for drawing graph, updates node limit if necessary
369 def drawGraph(self, n_limit):
370     self.n_limit = n_limit
371     self._drawNodesAndEdges()
372
373 # responsible for drawing nodes and edges in 3D space
374 def _drawNodesAndEdges(self):
375     # positions camera
376     camera(self.canvas_width/2, self.canvas_height/2, self.canvas_width*-2,
377 0.5,0.5,0, 0,1,0)
378
379     # updates rotation
380     self.rot = (self.rot[0], self.rot[1]-math.pi/100, self.rot[2])
381
382     background(0)
383     strokeWeight(2)
384     stroke(255)
385     fill(255)
386
387     for n in range(self.num_nodes):
388         pushMatrix()
389
390     # sets new rotation

```

```

386         rotateZ(self.rot[2])
387         rotateY(-1*self.rot[1])
388
389         # sets drawing origin to current node
390         translate((self.nodes[n][0])*self.canvas_width, (self.nodes[n][1])*
self.canvas_height, (self.nodes[n][2])*self.canvas_width)
391
392         # places ellipse at origin
393         ellipse(0, 0, 10, 10)
394
395         # draw all edges
396         if self.num_nodes <= self.n_limit:
397             for e_i in self.edges[self.nodes[n]]:
398                 e = self.nodes[e_i]
399                 # draws line from origin to neighboring node
400                 line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
- self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
canvas_width)
401             # draw edges for min degree node
402             elif self.nodes[n] == self.minDeg:
403                 stroke(0,255,0)
404                 for e_i in self.edges[self.nodes[n]]:
405                     e = self.nodes[e_i]
406                     # draws line from origin to neighboring node
407                     line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
- self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
canvas_width)
408                 stroke(255)
409             # draw edges for max degree node
410             elif self.nodes[n] == self.maxDeg:
411                 stroke(0,0,255)
412                 for e_i in self.edges[self.nodes[n]]:
413                     e = self.nodes[e_i]
414                     # draws line from origin to neighboring node
415                     line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
- self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
canvas_width)
416                 stroke(255)
417
418         popMatrix()

```