

Backbone Determination in a Wireless Sensor Network

Jake Carlson

February 18, 2018

Abstract

A report on implementing algorithms to partition a random geometric graph into bipartite subgraphs. Three different graph geometries are explored: unit square, unit disk, and unit sphere. Nodes are uniformly distributed in the geometry. Then the edges are determined and the vertices are colored using smallest-last vertex ordering.

Contents

1	Executive Summary	3
1.1	Introduction	3
1.2	Environment Description	3
2	Reduction to Practice	3
2.1	Data Structure Design	3
2.2	Algorithm Descriptions	4
2.2.1	Node Placement	4
2.2.2	Edge Determination	4
2.2.3	Graph Coloring	5
2.3	Algorithm Engineering	5
2.3.1	Node Placement	5
2.3.2	Edge Determination	6
2.3.3	Graph Coloring	6
2.4	Verification	7
2.4.1	Node Placement	7
2.4.2	Edge Determination	7
2.4.3	Graph Coloring	7
3	Appendix A - Figures	9
4	Appendix B - Code Listings	16

Listings

1	Processing driver	16
2	Topology class and subclasses	16

Benchmark	Order	A	Topology	r	Size	Realized A	Max Deg	Min Deg	Run Time (s)
1	1000	32	Square	0.101	14617	29	45	6	0.092
2	8000	64	Square	0.050	245861	61	98	15	1.248
3	16000	32	Square	0.025	250205	31	52	8	1.342
4	64000	64	Square	0.018	2014587	62	96	15	10.000
5	64000	128	Square	0.025	4009474	125	173	32	18.292
6	128000	64	Square	0.013	4053485	63	100	16	19.323
7	128000	128	Square	0.018	8068821	126	175	32	36.591
8	8000	64	Disk	0.045	246753	61	90	16	1.130
9	64000	64	Disk	0.016	2017701	63	98	18	9.627
10	64000	128	Disk	0.022	4017779	125	174	42	18.197
11	16000	64	Sphere	0.126	511826	63	93	32	19.922
12	32000	128	Sphere	0.126	2052907	128	177	85	80.309
13	64000	128	Sphere	0.089	4094853	127	168	79	145.549

Table 1: Benchmarks for Generating RGGs

1 Executive Summary

1.1 Introduction

Random geometric graphs (RGGs) are useful for simulating wireless sensor networks placed in different topologies. This project examines three different geometries: Square, Disk, and Sphere. The user supplies parameters for how many nodes they want in the network and how many connections they want for each node. Then, the simulation finds the average radius needed for that number of connections, finds multiple backbones for the RGG, and displays the results graphically.

1.2 Environment Description

The data structures and topologies for this simulation are implemented in Python2.7. The graphics are done using Processing.py. All development and benchmarking has been done on a 2014 MacBook Pro with a 3 GHz Intel Core i7 processor and 16 GB of DDR3 RAM running macOS High Sierra 10.13.3.

I elected to use Processing because of the simple API used to draw and render shapes in two- and three-dimensional space. I choose to use Processing.py over Java Processing because of my familiarity and comfort with the Python programming language. I could have written the graph generation, coloring, and backbone determination in another language and saved the data to a file for loading in processing, but since I intend to implement all of the algorithms in linear time, I don't think there will be a performance issue with using an interpreted language instead of a compiled one.

A separate data generation script was used to generate the summary table and graphs using matplotlib. This library, and a variety of others, could not be imported into Processing.py because the jython interpreter used by Processing only accepts libraries written in raw Python.

The different geometries were implemented in a stand alone Python file and imported into the Processing.py script or the data generation script depending on what was being run. These classes can then be used directly by Processing or the data generation script. Because there is no intermediary file to hold the generated nodes and edges, there is no additional disk space needed to run the simulation. Everything can be done in system memory managed by Processing.

2 Reduction to Practice

2.1 Data Structure Design

The primary data structure used for this project is an adjacency list. However, to allow for constant time lookup of edges of a node, I used a Python dictionary where the keys are nodes and the values are a list of indices of adjacent nodes in the original list of nodes. The space needed by the adjacency list is

Benchmark	Max Deg Deleted	Color Sets	Largest Color Set	Terminal Clique Size
1	20	20	76	18
2	40	39	323	35
3	24	23	1146	21
4	41	39	2543	34
5	73	66	1376	51
6	42	41	5049	36
7	74	66	2731	53
8	39	37	322	36
9	40	41	2547	40
10	73	66	1373	65
11	40	39	624	38
12	90	65	674	60
13	85	66	1351	60

Table 2: Benchmarks for Coloring RGGs

$\Theta(|V| + 2|E|)$. Two entries are used for each edge because they are undirected. This is superior to the adjacency matrix data structure which would require $\Theta(|E|^2)$ space.

In order to make this project maintainable as it is developed along the semester, I used the object-oriented capabilities Python has to offer to design the different geometries. I start with a base Topology class that creates the interface Processing uses to draw the graphs. This base class implements all of the methods needed for node placement and edge detection in 2D graphs. Then, I create three subclasses: Square, Disk, and Sphere.

The Square and Disk topologies simply need to override the methods for generating nodes and calculating the node radius needed for the desired average degree. The Sphere subclass needs to override a few additional functions because it exists in a 3D space. Other than the methods for generating nodes and calculating the node radius, it also needs to override the function used to draw the graph so that Processing will render the graph properly in 3D.

2.2 Algorithm Descriptions

2.2.1 Node Placement

A different node placement algorithm is required for each of the geometries. For the Square, the coordinates for each node are generated as two random numbers taken from a uniform distribution on the range $[0, 1]$. All of these points are guaranteed to be in the unit square.

For the Disk, a similar method is used. The coordinates for nodes are randomly sampled from a uniform distribution; however, if a node has a distance from the center of the Disk greater than the radius of 1, the coordinates for that node are resampled.

For the Sphere a different method must be used so that all of the nodes are placed on the surface of the Sphere and the volume is vacant. For this geometry, I used the following equations:

$$x = \sqrt{1 - u^2} \cos \theta \quad (1)$$

$$y = \sqrt{1 - u^2} \sin \theta \quad (2)$$

$$z = u \quad (3)$$

where $\theta \in [0, 2\pi]$ and $u \in [-1, 1]$. This is guaranteed to uniformly distribute nodes on the surface area of the sphere [1].

All of these algorithms can be solved in $\Theta(|V|)$ where because each node only needs to be assigned a position once.

2.2.2 Edge Determination

There are several methods for finding the edges in the graph. The brute force method checks every node, and for each node checks all other nodes to see if they are close enough to form an edge. The brute force method is $\Theta(|V|^2)$.

The second method to find the edges is the sweep method. This method first sorts the nodes along the x-axis. Then, for any node, we only need to search left and right until the distance along the x-axis is greater than the connection radius for the nodes. This dramatically reduces the search space. The sweep method is $O(n \lg(n) + 2rn^2)$ where $n = |V|$ and r is the connection radius. The $n \lg(n)$ portion is for the sorting and the $2rn^2$ portion is for measuring the distance between nodes in a sweep step.

The final method to find edges is the cell method. This method places the nodes into cells of area $r \times r$ based on their position in the topology. When the edge detection runs, each node needs to be visited once, but only the cell the node populates and the neighboring cells need to be searched for connections.

The only method that needs to be adjusted for the Sphere is the cell method. Instead of using a two dimensional grid of cells, a three dimensional mesh is needed to divide the topology. The cells then have volume $r \times r \times r$. Only the current cell and the neighboring cells need to be searched.

2.2.3 Graph Coloring

Two algorithms are used for coloring the graphs. The first is smallest-last vertex ordering, which sorts the vertices based on the number of degrees they have. The second is the greedy graph coloring algorithm.

Smallest-last vertex ordering is used to order the nodes for coloring. The steps to this algorithm are as follows [3]:

1. Initialize a representation of your target graph
2. Find the vertex v_j of minimum degree in your representation
3. Update your representation to simulate deleting v_j
4. If there are still vertices in the representation, return to step 1, otherwise terminate with the sequence of vertices removed

This algorithm is linear if each of the above steps is linear. Step 1 is linear if we can build a representation of the graph in linear time. For this, we can use an array of buckets, where each bucket holds the vertices that have the same number of edges as the position of the bucket in the array of buckets. To build this data structure, each node only needs to be visited once, making this linear in both space and time. Next, finding the vertex of minimum degree simply requires finding the lowest index bucket that has a node. This is bounded by the number of buckets, which is bounded by the number of nodes, making Step 2 linear. Next, we have to update the representation of the graph. To do this, we have to look at each node that shares an edge with v_j and move it to the bucket for nodes with one fewer degree. This requires traversing the list of edges for v_j which means Step 3 is linear. Since this is repeated for each node, the runtime of this program is $\Theta(|E| + |V|)$ and the space needed is $\Theta(|V|)$.

After this, a single traversal of the smallest-last vertex ordering is needed to color the graph. As we traverse this list, we check to see if the nodes before it (that are already colored) share an edge with the current node. The node can then be colored with any color it does not share an edge with or, if it shares an edge with all currently used colors, it is assigned a new color. This algorithm is also linear. Each node needs to be visited once and when a node is visited, all previous nodes are checked to see if they are in the edge list of the current node. Because we used smallest last vertex ordering, as we have to check more and more nodes, we get to check fewer and fewer edges. This makes the greedy coloring algorithm $O(|V| + |E|)$.

2.3 Algorithm Engineering

2.3.1 Node Placement

It is easy to implement the algorithms for placing nodes in the different geometries using Python's math library. This library offers functions for sampling points on a uniform distribution. For the Square, sampling on a range $[0, 1]$ is sufficient for all of the nodes. Since each node only needs to be placed once, this runs at $\Theta(|V|)$ where:

For the Disk, the node needs to be resampled if it is too far from the center. To do this, the distance function is used to find the distance between the node and the center. If the node is further than 1 from the center, node generation falls into a while loop which iterates until the node is within the unit circle. Since nodes are taken from a uniform distribution, the number of nodes that will need to be resampled is approximately equal to the ratio of the area of the square that circumscribes the unit circle which falls outside of the unit circle to the total area of the square. This is given by:

$$\frac{(2r)^2 - \pi r^2}{(2r)^2} = \frac{4 - \pi}{4} = 0.2146 \quad (4)$$

Since the placement algorithm for each node of the Disk will iterate until the node falls within the unit circle, the total number of iterations N can be found as the sum of the geometric series:

$$N = \sum_{k=0}^{\infty} n(0.2146)^k = \frac{n}{1 - 0.2146} = 1.273n \quad (5)$$

where $n = |V|$. This shows this implementation is $\Theta(n)$.

For the node placement algorithm of the Sphere, again the math library in Python makes this easy. Each node needs two random values pulled from a uniform distribution, two square root operations, one sine operation, and one cosine operation. Each node only needs to be placed once so the runtime of this algorithm is $\Theta(n)$ where $n = |V|$.

2.3.2 Edge Determination

Each method implemented for finding edges has a different time complexity. The brute force method uses an outer loop and an inner loop, which each iterate over every node in the graph. An edge is saved to the adjacency list if the nodes are not the same and the distance between them is less than or equal to the calculated node radius. This is guaranteed to run in $\Theta(n^2)$ where $n = |V|$. The number of times the distance needs to be calculated is $n \times (n - 1)$ because it will not be calculated when the nodes are the same (distance would be zero, but no edge is drawn here). No additional space is needed for the brute force method so the space complexity is $O(1)$.

The implementation of sweep starts by sorting the nodes along the x-axis. Python lists have a built-in sort function that has $O(n \lg(n))$ time complexity [2]. After this stage, it iterates over every node building a search space which will be scanned for edges. For each node, the list of nodes is searched right $r \times n$ nodes to find those within one radius length of the current node. With the search space built, the search space is iterated over once to find nodes that have a distance less than or equal the node radius. Then, the indices of the nodes are added to the adjacency list entry for each other. My implementation of this runs in $O(n \lg(n) + 2rn)$ where $n = |V|$ and r is the node connection radius. Because the list sort method sorts inplace, the only additional space needed is for the search space. This saves $O(rn)$ nodes and is reset after every iteration.

The cell method implementation works in linear time. In the first step of the method, the cells are initialized as a list of empty lists. There are $(1/r + 1)^2$ cells. The nodes are then iterated over and assigned a cell by dividing their x and y coordinates by the node radius. At this point, the cells are iterated over and, for each node in the cell, the nodes in the current cell and the four forward adjacent cells and the are checked to see if they fall within the node radius of the current node. All together, this implementation runs at $O(n + n + 5nr^2) = O((2 + 5r^2)n)$ where $n = |V|$. The amount of additional space needed is equal to the number of nodes because they are copied into their respective cells. This places the space complexity at $\Theta(n)$.

2.3.3 Graph Coloring

Implementing the smallest-last coloring algorithm involves implementing the smallest-last vertex ordering algorithm and the greedy graph coloring algorithm. For smallest-last vertex ordering, the first thing to do is build the data structure used to represent the graph with deleted nodes. The number of sets needed is equal to the maximum degree of the nodes. Then, the index of each node is placed in the set corresponding to the number of edges it has then the RGG. Simultaneously, a dictionary is created that maps each node to the number of degrees it has in the graph with deletions. Each value starts at the number of edges the corresponding node has in the RGG. At this point, we have iterated over all of the nodes once and allocated space for twice the number of nodes by copying them into the sets and using them as the keys for the degrees dictionary.

Because Python dictionaries resize at specific numbers of entries, we can determine the number of additional insertions caused by rehashing while the degrees dictionary is built. Python dictionaries start out with space for 8 entries and quadruple in size until the number of entries is above 50,000, at which point it begins to double in size. Clearly the dictionary grows at a logarithmic rate, but the total number of insertions I for an input size of n is given by:

$$I = \begin{cases} n + 8 \sum_{k=1}^{\log_4 \lceil n/8 \rceil} 4^k & n \leq 50,000 \\ n + 8 \sum_{k=1}^6 4^k + 32768 \sum_{k=1}^{\log_2 \lceil n/32768 \rceil} 2^k & n > 50,000 \end{cases} \quad (6)$$

Fortunately, because the entire dictionary is built before it is used by the smallest-last vertex ordering algorithm, it will never again be resized once the algorithm starts. Unfortunately, the sets resize at a similar rate and it is more difficult to predict how large the sets will need to be when performing smallest-last vertex ordering. The degree dictionary will also be used to index into the sets, so we gain a speed up here by not having to iterate over all of the edges for a node and determining if the node it shares an edge with are in the remaining graph each time we want to sift nodes down to lower set.

Next, the smallest-last vertex ordering algorithm is run until every node has been removed from the sets. For each node, I iterate over the sets from lowest degree to highest degree to find the first non-empty set. This set must contain the next node to remove because it contains all nodes with smallest degree. Before deleting the node from the graph and moving all adjacent nodes down a set, I check to see if the current set has all remaining nodes. If this is the case, the terminal clique has been found, and the size of the terminal clique must be saved. After this check, a node is popped from the end of the current set, and appended to the smallest-last ordering result. Then, for all the adjacent nodes to the popped node in the original graph, I check if the node is in the set with its degree. If it is, the number of degrees for that node can be decremented and the node can be placed into the correct set for its new degree.

The last step is to reverse the order of the smallest-last ordering result because it was built in the opposite order (smallest-first). All together, excluding the initialization of accessory data structures, this implementation runs in $\Theta(2|V| + 2|E|)$ time and $\Theta(2|V|)$ space since nodes are removed from the buckets and added to the result.

After this the graph needs to be colored. For this I initially assign each node a color of -1 in a node color array that is parallel to the original list of nodes. I iterate over all of the nodes in the smallest-last vertex ordering. At each node, I generate a set of colors that is already used by the neighbors of that node by iterating over all of its edge nodes and grabbing their color from the node color array. Then, I just have to increment color from 0 until it does not exist in the search space set and I have the color to assign to the node.

Since the smallest-last ordering is used, each time I check to see if a node is adjacent to the current node, I am searching nodes with fewer and fewer edges. This means that the nodes with the most neighbors are searched first, when the number of other nodes to check is lowest, and the nodes with the fewest neighbors are searched last, when we have the most nodes to check if they share an edge with the current node. All together, this implementation runs in $\Theta(|V| + 2|E|)$ time and $\Theta(|V|)$ space because we need a new array for the colors assigned to each of the nodes.

2.4 Verification

2.4.1 Node Placement

The nodes can be verified to be distributed uniformly if the distribution of degrees follows a normal distribution. To show that the distribution of degrees for each of my geometries are following a normal distribution, I plotted degree histograms for each of the geometries with 32,000 nodes and an average degree of 16. The histogram for Square is given in Figure 1, Disk is given in Figure 2, and Sphere is given in Figure 3. These histograms clearly follow a normal distribution.

2.4.2 Edge Determination

The runtime for the edge detection methods can be verified by varying the number of nodes and measuring the runtime of each algorithm. By looking at how the runtime grows, we can calculate the trendline that best fits the growth rate. For the first comparison, I vary the number of nodes from 4,000 to 64,000 in steps of 4,000, while holding the desired average degree constant at 16. As we can see in Figure 4, the growth rates of the brute force and sweep methods are quadratic, while the growth rate of the cell method. The trendline functions are given on the graph.

For the second metric, I held the number of nodes constant at 32,000 and varied the desired average degree from 2 to 32 in steps of 2. The graph is given in Figure 5. The cell method clearly grows linearly, but the sweep method is harder to gauge. Since varying the desired average degree should only change the node radius, I would expect this to grow linearly as well. However, because each graph is randomly generated, some graphs can have nodes that are closer to sorted order than others. This can effect the measured runtime. It would be easier to gauge the trend if it I ran the data collection multiple times and averaged the results.

2.4.3 Graph Coloring

References

- [1] Weisstein, Eric W., Wolfram MathWorld
Sphere Point Picking
<http://mathworld.wolfram.com/SpherePointPicking.html>
- [2] Peters, Tim
Timsort
<http://svn.python.org/projects/python/trunk/Objects/listsort.txt>
- [3] Matula, David and Beck, Leland
Smallest-Last Ordering and Clustering and Graph Coloring Algorithms
- [4] Johnson, Ian
Linear-Time Computation of High-Converage Backbones for Wireless Sensor Networks
<https://github.com/ianjohnson/SensorNetwork/blob/master/Report/Report.pdf>
- [5] Rees, Gareth *Python's underlying hash data structure for dictionaries*
<https://stackoverflow.com/questions/4279358/pythons-underlying-hash-data-structure-for-dictionaries>
- [6] Thomas, Alec *Why is tuple faster than list?* <https://stackoverflow.com/questions/3340539/why-is-tuple-faster-than>
- [7] Kruse, Lars *Python Speed, Performance Tips* <https://wiki.python.org/moin/PythonSpeed/PerformanceTips>

3 Appendix A - Figures

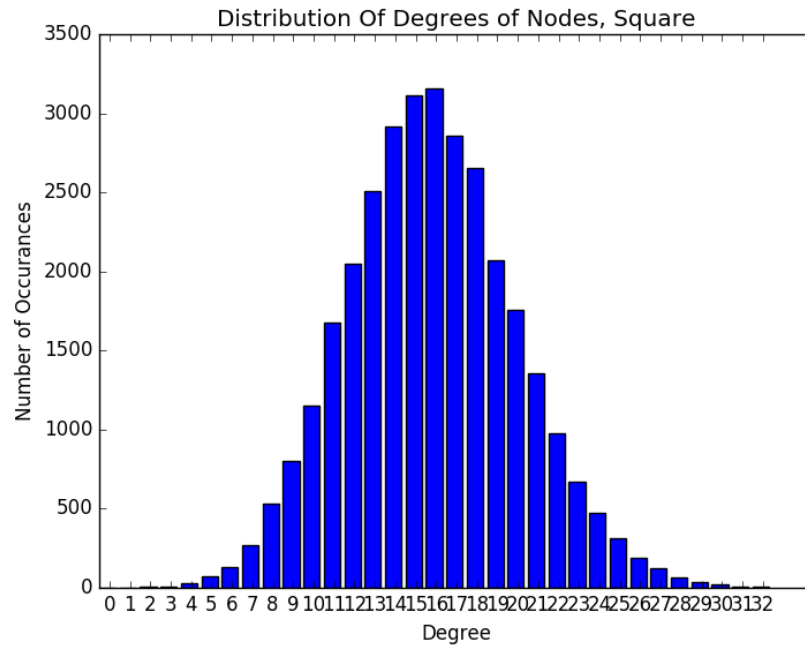


Figure 1: Distribution of Degree counts for Square. 32,000 Nodes, Average Degree of 16

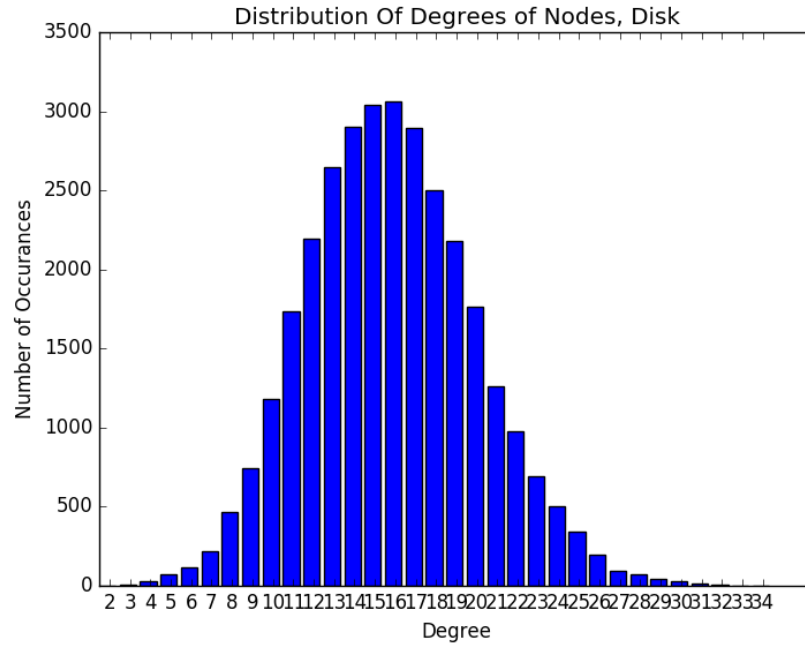


Figure 2: Distribution of Degree counts for Disk. 32,000 Nodes, Average Degree of 16

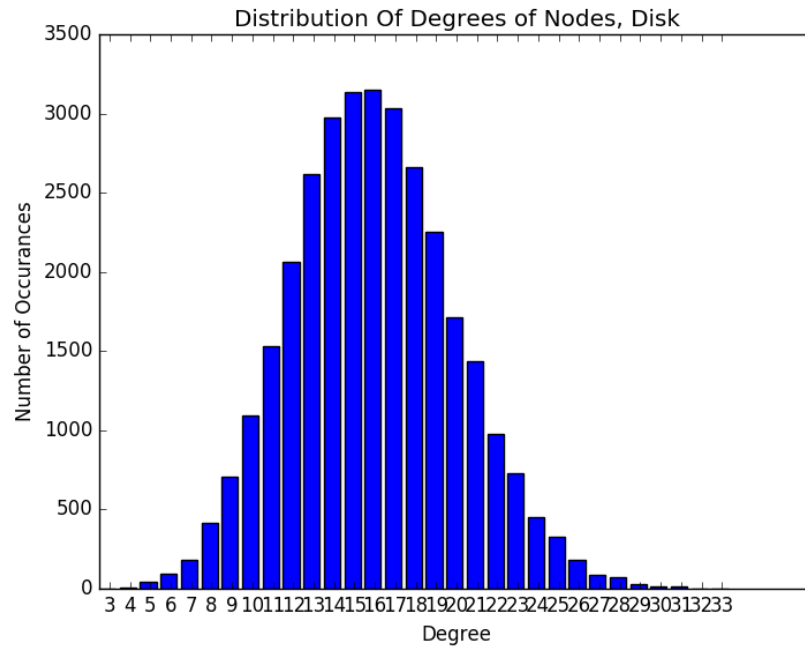


Figure 3: Distribution of Degree counts for Sphere. 32,000 Nodes, Average Degree of 16

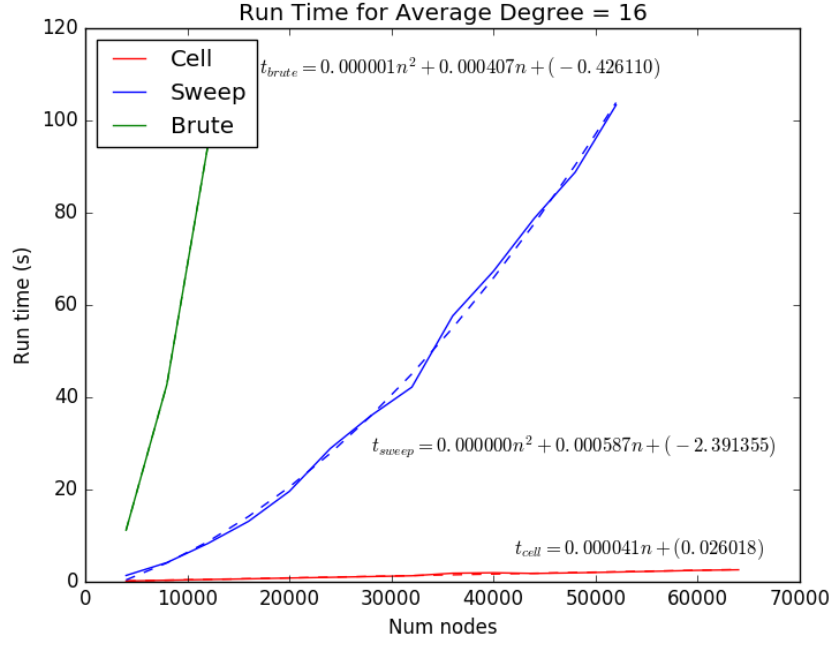


Figure 4: Runtime for Each Edge Detection Method, Average Degree of 16

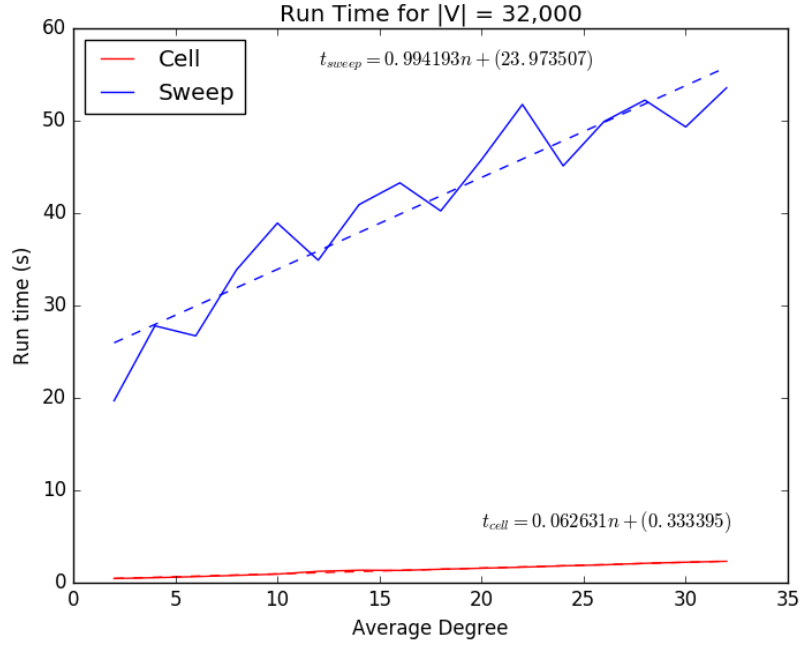


Figure 5: Runtime for Cell and Sweep Edge Detection, Variable Average Degree

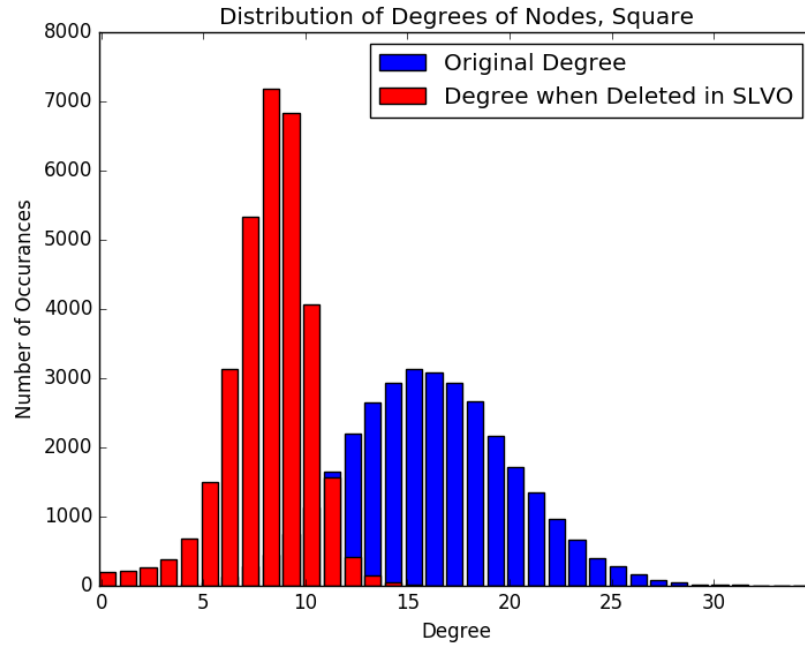


Figure 6: Distribution of Degree when Deleted for Square. 32,000 Nodes, Average Degree of 16

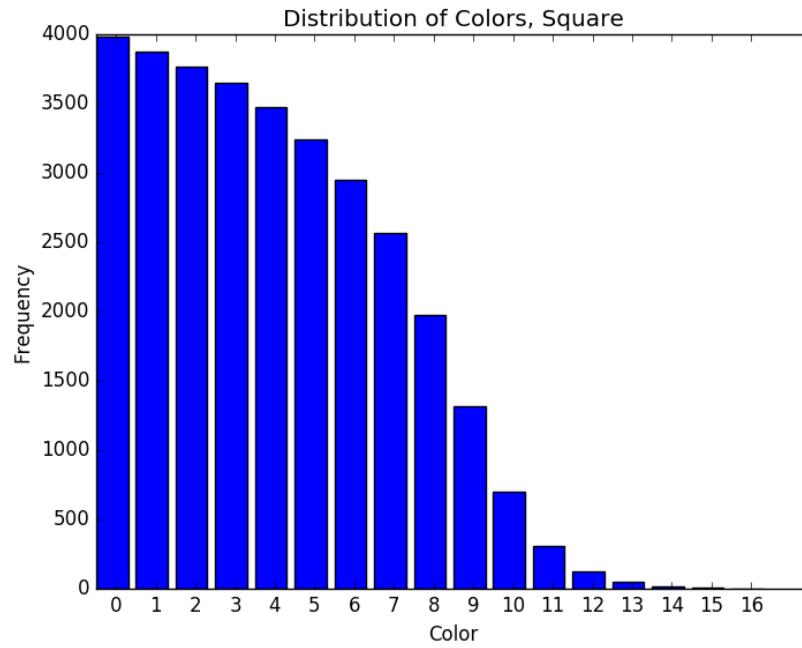


Figure 7: Distribution of Colors for Square. 32,000 Nodes, Average Degree of 16




Figure 8: Square benchmark graphs






Figure 9: Disk benchmark graphs






Figure 10: Sphere benchmark graphs

4 Appendix B - Code Listings

Listing 1: Processing driver

```
1 import random
2 import time
3 import math
4 from collections import Counter
5 from objects.topology import Square, Disk, Sphere
6
7 CANVAS_HEIGHT = 720
8 CANVAS_WIDTH = 720
9
10 NUMNODES = 100
11 AVG_DEG = 10
12
13 MAX_NODES_TO_DRAW_EDGES = 8000
14
15 RUN_BENCHMARK = False
16
17 def setup():
18     size(CANVAS_WIDTH, CANVAS_HEIGHT, P3D)
19     background(0)
20
21 def draw():
22     topology.drawGraph(MAX_NODES_TO_DRAW_EDGES)
23
24 def main():
25     global topology
26     # topology = Square()
27     # topology = Disk()
28     topology = Sphere()
29
30     topology.num_nodes = NUMNODES
31     topology.avg_deg = AVG_DEG
32     topology.canvas_height = CANVAS_HEIGHT
33     topology.canvas_width = CANVAS_WIDTH
34
35     if RUN_BENCHMARK:
36         n_benchmark = 0
37         topology.prepBenchmark(n_benchmark)
38
39     run_time = time.clock()
40
41     topology.generateNodes()
42     topology.findEdges(method="cell")
43     topology.colorGraph()
44
45     print "Average degree: {}".format(topology.findAvgDegree())
46     print "Min degree: {}".format(topology.getMinDegree())
47     print "Max degree: {}".format(topology.getMaxDegree())
48     print "Num edges: {}".format(topology.findNumEdges())
49     print "Terminal clique size: {}".format(topology.term_clique_size)
50     print "Number of colors: {}".format(len(set(topology.node_colors)))
51     print "Max degree when deleted: {}".format(max(topology.deg_when_del.values()))
52
53     color_cnt = Counter(topology.node_colors)
54     print "Max color set size: {} color: {}".format(color_cnt.most_common(1)
55     [0][1],
56     color_cnt.most_common(1)
57     [0][0])
58
59     run_time = time.clock() - run_time
60     print "Run time: {:.3f} s".format(run_time)
61
62 main()
```


Listing 2: Topology class and subclasses

```

1 import random
2 import math
3 import time
4
5 # benchmarks (num_nodes, avg_deg)
6 SQUARE_BENCHMARKS = [(1000,32), (8000,64), (16000,32), (64000,64), (64000,128),
7                       (128000,64), (128000, 128)]
8 DISK_BENCHMARKS = [(8000,64), (64000,64), (64000,128)]
9 SPHERE_BENCHMARKS = [(16000,64), (32000,128), (64000,128)]
10
11 """
12 Topology – super class for the shape of the random geometric graph
13 """
14 class Topology(object):
15
16     num_nodes = 100
17     avg_deg = 0
18     canvas_height = 720
19     canvas_width = 720
20
21     def __init__(self):
22         self.nodes = []
23         self.edges = {}
24         self.node_r = 0.0
25         self.minDeg = ()
26         self.maxDeg = ()
27         self.s_last = []
28         self.deg_when_del = {}
29         self.node_colors = []
30
31     # public function for generating nodes of the graph, must be subclassed
32     def generateNodes(self):
33         print "Method for generating nodes not subclassed"
34
35     # public function for finding edges
36     def findEdges(self, method="brute"):
37         self._getRadiusForAverageDegree()
38         self._addNodesAsEdgeKeys()
39
40         if method == "brute":
41             self._bruteForceFindEdges()
42         elif method == "sweep":
43             self._sweepFindEdges()
44         elif method == "cell":
45             self._cellFindEdges()
46         else:
47             print "Find edges method not defined: {}".format(method)
48
49         self._findMinAndMaxDegree()
50
51     # brute force edge detection
52     def _bruteForceFindEdges(self):
53         for i, n in enumerate(self.nodes):
54             for j, m in enumerate(self.nodes):
55                 if i != j and self._distance(n, m) <= self.node_r:
56                     self.edges[n].append(j)
57
58     # sweep edge detection
59     def _sweepFindEdges(self):
60         self.nodes.sort(key=lambda x: x[0])
61
62         for i, n in enumerate(self.nodes):
63             search_space = []
64             for j in range(1, self.num_nodes-i):
65                 if abs(n[0] - self.nodes[i+j][0]) <= self.node_r:
66                     search_space.append(i+j)
67             else:

```

```

68         break
69     for j in search_space:
70         if self._distance(n, self.nodes[j]) <= self.node_r:
71             self.edges[n].append(j)
72             self.edges[self.nodes[j]].append(i)
73
74 # cell edge detection
75 def _cellFindEdges(self):
76     num_cells = int(1/self.node_r) + 1
77     cells = []
78     for i in range(num_cells):
79         cells.append([[j for j in range(num_cells)]])
80
81     for i, n in enumerate(self.nodes):
82         cells[int(n[0]/self.node_r)][int(n[1]/self.node_r)].append(i)
83
84     for i in range(num_cells):
85         for j in range(num_cells):
86             for n_i in cells[i][j]:
87                 for c in self._findAdjCells(i, j, num_cells):
88                     for m_i in cells[c[0]][c[1]]:
89                         if self._distance(self.nodes[n_i], self.nodes[m_i]) <=
self.node_r:
90                             self.edges[self.nodes[n_i]].append(m_i)
91                             self.edges[self.nodes[m_i]].append(n_i)
92                         for m_i in cells[i][j]:
93                             if self._distance(self.nodes[n_i], self.nodes[m_i]) <=
self.node_r and n_i != m_i:
94                                 self.edges[self.nodes[n_i]].append(m_i)
95
96 # cell edge detection helper function
97 def _findAdjCells(self, i, j, n):
98     adj_cells = [(1,-1), (0,1), (1,1), (1,0)]
99     return (((i+x[0])%n, (j+x[1])%n) for x in adj_cells)
100
101 # function for finding the radius needed for the desired average degree
102 # must be subclassed
103 def _getRadiusForAverageDegree(self):
104     print "Method for finding necessary radius for average degree not
subclassed"
105
106 # helper function for findEdges, initializes edges dict
107 def _addNodesAsEdgeKeys(self):
108     self.edges = {n:[] for n in self.nodes}
109
110 # calculates the distance between two nodes (2D)
111 def _distance(self, n, m):
112     return math.sqrt((n[0] - m[0])**2 + (n[1] - m[1])**2)
113
114 # public function for finding the number of edges
115 def findNumEdges(self):
116     sigma_edges = 0
117     for k in self.edges.keys():
118         sigma_edges += len(self.edges[k])
119
120     return sigma_edges/2
121
122 # public function for finding the average degree of nodes
123 def findAvgDegree(self):
124     return 2*self.findNumEdges()/self.num_nodes
125
126 # helper function for finding nodes with min and max degree
127 def _findMinAndMaxDegree(self):
128     self.minDeg = self.edges.keys()[0]
129     self.maxDeg = self.edges.keys()[0]
130
131     for k in self.edges.keys():
132         if len(self.edges[k]) < len(self.edges[self.minDeg]):

```

```

133         self.minDeg = k
134         if len(self.edges[k]) > len(self.edges[self.maxDeg]):
135             self.maxDeg = k
136
137     # public function for getting the minimum degree
138     def getMinDegree(self):
139         return len(self.edges[self.minDeg])
140
141     # public function for getting the maximum degree
142     def getMaxDegree(self):
143         return len(self.edges[self.maxDeg])
144
145     # public function for setting up the benchmark to run, must be subclassed
146     def prepBenchmark(self, n):
147         print "Method for preparing benchmark not subclassed"
148
149     # public function for drawing the graph
150     def drawGraph(self, n_limit):
151         self._drawNodes()
152         if self.num_nodes <= n_limit:
153             self._drawEdges()
154         else:
155             self._drawMinMaxDegNodes()
156
157     # responsible for drawing the nodes in the canvas
158     def _drawNodes(self):
159         strokeWeight(2)
160         stroke(255)
161         fill(255)
162
163         for n in range(self.num_nodes):
164             ellipse(self.nodes[n][0]*self.canvas_width, self.nodes[n][1]*self.
canvas_height, 5, 5)
165
166     # responsible for drawing the edges in the canvas
167     def _drawEdges(self):
168         strokeWeight(1)
169         stroke(245)
170         fill(255)
171
172         for n in self.edges.keys():
173             for m_i in self.edges[n]:
174                 line(n[0]*self.canvas_width, n[1]*self.canvas_height, self.nodes[
m_i][0]*self.canvas_width, self.nodes[m_i][1]*self.canvas_height)
175
176     # responsible for drawing the edges of the min and max degree nodes
177     def _drawMinMaxDegNodes(self):
178         strokeWeight(1)
179         stroke(0,255,0)
180         fill(255)
181         for n_i in self.edges[self.minDeg]:
182             line(self.minDeg[0]*self.canvas_width, self.minDeg[1]*self.
canvas_height, self.nodes[n_i][0]*self.canvas_width, self.nodes[n_i][1]*self.
canvas_height)
183
184         stroke(0,0,255)
185         for n_i in self.edges[self.maxDeg]:
186             line(self.maxDeg[0]*self.canvas_width, self.maxDeg[1]*self.
canvas_height, self.nodes[n_i][0]*self.canvas_width, self.nodes[n_i][1]*self.
canvas_height)
187
188     # uses smallest last vertex ordering to color the graph
189     def colorGraph(self):
190         self.s_last, self.deg_when_del = self._smallestLastVertexOrdering()
191         self.node_colors = self._assignNodeColors(self.s_last)
192
193     # constructs a degree structure and determines the smallest last vertex
ordering

```

```

194     def _smallestLastVertexOrdering(self):
195         deg_sets = {l:set() for l in range(len(self.edges[self.maxDeg])+1)}
196         deg_when_del = {n:len(self.edges[n]) for n in self.nodes}
197
198         for i, n in enumerate(self.nodes):
199             deg_sets[deg_when_del[n]].add(i)
200
201         smallest_last_ordering = []
202
203         clique_found = False
204         j = len(self.nodes)
205         while j > 0:
206             # get the current smallest bucket
207             curr_bucket = 0
208             while len(deg_sets[curr_bucket]) == 0:
209                 curr_bucket += 1
210
211             # if all the remaining nodes are connected we have the terminal clique
212             if not clique_found and len(deg_sets[curr_bucket]) == j:
213                 clique_found = True
214                 self.term_clique_size = curr_bucket
215
216             # get node with smallest degree
217             v_i = deg_sets[curr_bucket].pop()
218             smallest_last_ordering.append(v_i)
219
220             # decrement position of nodes that shared an edge with v
221             for n_i in (n_i for n_i in self.edges[self.nodes[v_i]] if n_i in
deg_sets[deg_when_del[self.nodes[n_i]]]):
222                 deg_sets[deg_when_del[self.nodes[n_i]]].remove(n_i)
223                 deg_when_del[self.nodes[n_i]] -= 1
224                 deg_sets[deg_when_del[self.nodes[n_i]]].add(n_i)
225
226             j -= 1
227
228             # reverse list since it was built shortest-first
229             return smallest_last_ordering[::-1], deg_when_del
230
231     # assigns the colors to nodes given in a smallest-last vertex ordering as a
parallel array
232     def _assignNodeColors(self, s_last):
233         colors = [-1 for _ in range(len(s_last))]
234         for i in s_last:
235             adj_colors = set([colors[j] for j in self.edges[self.nodes[i]]])
236             color = 0
237             while color in adj_colors:
238                 color += 1
239             colors[i] = color
240
241         return colors
242
243     """
244     Square – inherits from Topology, overloads generateNodes and
_getRadiusForAverageDegree
245     for a unit square topology
246     """
247     class Square(Topology):
248
249         def __init__(self):
250             super(Square, self).__init__()
251
252         # places nodes uniformly in a unit square
253         def generateNodes(self):
254             for i in range(self.num_nodes):
255                 self.nodes.append((random.uniform(0,1), random.uniform(0,1)))
256
257         # calculates the radius needed for the requested average degree in a unit
square

```

```

258     def _getRadiusForAverageDegree(self):
259         self.node_r = math.sqrt(self.avg_deg/(self.num_nodes * math.pi))
260
261     # gets benchmark setting for square
262     def prepBenchmark(self, n):
263         self.num_nodes = SQUARE_BENCHMARKS[n][0]
264         self.avg_deg = SQUARE_BENCHMARKS[n][1]
265
266     """
267     Disk – inherits from Topology, overloads generateNodes and
268         _getRadiusForAverageDegree
269     for a unit circle topology
270     """
271     class Disk(Topology):
272         def __init__(self):
273             super(Disk, self).__init__()
274
275         # places nodes uniformly in a unit square and regenerates the node if it falls
276         # outside of the circle
277         def generateNodes(self):
278             for i in range(self.num_nodes):
279                 p = (random.uniform(0,1), random.uniform(0,1))
280                 while self._distance(p, (0.5,0.5)) > 0.5:
281                     p = (random.uniform(0,1), random.uniform(0,1))
282                 self.nodes.append(p)
283
284         # calculates the radius needed for the requested average degree in a unit
285         # circle
286         def _getRadiusForAverageDegree(self):
287             self.node_r = math.sqrt((self.avg_deg + 0.0)/self.num_nodes)/2
288
289         # gets benchmark setting for disk
290         def prepBenchmark(self, n):
291             self.num_nodes = DISK_BENCHMARKS[n][0]
292             self.avg_deg = DISK_BENCHMARKS[n][1]
293
294     """
295     Sphere – inherits from Topology, overloads generateNodes,
296         _getRadiusForAverageDegree,
297         and _distance for a unit sphere topology. Also updates the drawGraph function for
298         a 3D canvas
299     """
300     class Sphere(Topology):
301         # adds rotation and node limit variables
302         def __init__(self):
303             super(Sphere, self).__init__()
304             self.rot = (0, math.pi/4, 0) # this may move to Topology if rotation is
305             # given to the 2D shapes
306             # used to control _drawNodes functionality
307             self.n_limit = 8000
308
309         # places nodes in a unit cube and projects them onto the surface of the sphere
310         def generateNodes(self):
311             for i in range(self.num_nodes):
312                 # equations for uniformly distributing nodes on the surface area of
313                 # a sphere: http://mathworld.wolfram.com/SpherePointPicking.html
314                 u = random.uniform(-1,1)
315                 theta = random.uniform(0, 2*math.pi)
316                 p = (
317                     math.sqrt(1 - u**2) * math.cos(theta),
318                     math.sqrt(1 - u**2) * math.sin(theta),
319                     u
320                 )
321                 self.nodes.append(p)
322
323         # calculates the radius needed for the requested average degree in a unit

```

```

322 sphere
323 def _getRadiusForAverageDegree(self):
324     self.node_r = math.sqrt((self.avg_deg + 0.0)/self.num_nodes)*2
325
326 # calculates the distance between two nodes (3D)
327 def _distance(self, n, m):
328     return math.sqrt((n[0] - m[0])**2+(n[1] - m[1])**2+(n[2] - m[2])**2)
329
330 # gets benchmark setting for sphere
331 def prepBenchmark(self, n):
332     self.num_nodes = SPHERE.BENCHMARKS[n][0]
333     self.avg_deg = SPHERE.BENCHMARKS[n][1]
334
335 # public function for drawing graph, updates node limit if necessary
336 def drawGraph(self, n_limit):
337     self.n_limit = n_limit
338     self._drawNodesAndEdges()
339
340 # responsible for drawing nodes and edges in 3D space
341 def _drawNodesAndEdges(self):
342     # positions camera
343     camera(self.canvas_width/2, self.canvas_height/2, self.canvas_width*-2,
344            0.5,0.5,0, 0,1,0)
345
346     # updates rotation
347     self.rot = (self.rot[0], self.rot[1]-math.pi/100, self.rot[2])
348
349     background(0)
350     strokeWeight(2)
351     stroke(255)
352     fill(255)
353
354     for n in range(self.num_nodes):
355         pushMatrix()
356
357         # sets new rotation
358         rotateZ(self.rot[2])
359         rotateY(-1*self.rot[1])
360
361         # sets drawing origin to current node
362         translate((self.nodes[n][0])*self.canvas_width, (self.nodes[n][1])*
363                self.canvas_height, (self.nodes[n][2])*self.canvas_width)
364
365         # places ellipse at origin
366         ellipse(0, 0, 10, 10)
367
368         # draw all edges
369         if self.num_nodes <= self.n_limit:
370             for e_i in self.edges[self.nodes[n]]:
371                 e = self.nodes[e_i]
372                 # draws line from origin to neighboring node
373                 line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
374                - self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
375                canvas_width)
376
377             # draw edges for min degree node
378             elif self.nodes[n] == self.minDeg:
379                 stroke(0,255,0)
380                 for e_i in self.edges[self.nodes[n]]:
381                     e = self.nodes[e_i]
382                     # draws line from origin to neighboring node
383                     line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
384                    - self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
385                    canvas_width)
386
387                 stroke(255)
388             # draw edges for max degree node
389             elif self.nodes[n] == self.maxDeg:
390                 stroke(0,0,255)
391                 for e_i in self.edges[self.nodes[n]]:

```

```

383         e = self.nodes[e_i]
384         # draws line from origin to neighboring node
385         line(0,0,0, (e[0] - self.nodes[n][0])*self.canvas_width, (e[1]
- self.nodes[n][1])*self.canvas_height, (e[2] - self.nodes[n][2])*self.
canvas_width)
386         stroke(255)
387
388         popMatrix()

```